

segalgouldn Directory cleaning and file exporting.

adb4c40 a minute ago

1 contributor

2177 lines (1655 sloc) | 53.2 KB

Using Scikit-Learn and NLTK to build a Naive Bayes Classifier that identifies subtweets

In all tables, assume:

- "❶" represents a single hashtag
- "❷" represents a single URL
- "❸" represents a single mention of username (e.g. "@noah")

Import libraries

```
%matplotlib inline
```

```
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.feature_extraction import text
from sklearn.naive_bayes import MultinomialNB
from sklearn.model_selection import KFold
from sklearn.pipeline import Pipeline
from sklearn.externals import joblib
from os.path import basename, splitext
from random import choice, sample
from nltk.corpus import stopwords
from string import punctuation
from pprint import pprint
from glob import glob

import matplotlib.pyplot as plt
import pandas as pd
import numpy as np

import scipy.stats
import itertools
import enchant
import nltk
import json
import re
```

Set up some regex patterns

```
hashtags_pattern = re.compile(r'(\#[a-zA-Z0-9]+)')
```

```
urls_pattern = re.compile(r'(?i)\b(?:https?://|www\d{0,3}[.]|[a-z0-9.-]+[.][a-z]{2,4}/)(?:[^\s()<>]|\\([^\s()<>]+|
```

```
at_mentions_pattern = re.compile(r'(<=^|(<=[^a-zA-Z0-9-\.]))@([A-Za-z0-9_]+)')
```

Prepare English dictionary for language detection

```
english_dict = enchant.Dict("en_US")
```

Use NLTK's tokenizer instead of Scikit's

```
tokenizer = nltk.casual.TweetTokenizer(preserve_case=False, reduce_len=True)
```

Prepare for viewing long text in CSVs and ones with really big and small numbers

```
pd.set_option("max_colwidth", 1000)
```

```
pd.options.display.float_format = "{:,.4f}".format
```

Load the two data files

Only use tweets with at least 50% English words

Also, make the mentions of usernames, URLs, and hashtags generic

```
def load_data(filename, threshold=0.5):
    data = [(hashtags_pattern.sub("ⓧ",
        urls_pattern.sub("ⓧ",
            at_mentions_pattern.sub("ⓧ",
                t["tweet_data"]["full_text"])))
        .replace("\u2018", "'')
        .replace("\u2019", "'')
        .replace("\u201c", "\"")
        .replace("\u201d", "\"")
        .replace("&quot;", "\"")
        .replace("&", "&")
        .replace(">", ">")
        .replace("<", "<"))
        for t in json.load(open(filename))
        if t["tweet_data"]["user"]["lang"] == "en"
        and t["reply"]["user"]["lang"] == "en"]
    new_data = []
    for tweet in data:
        tokens = tokenizer.tokenize(tweet)
        english_tokens = [english_dict.check(token) for token in tokens]
        percent_english_words = sum(english_tokens)/len(english_tokens)
        if percent_english_words >= threshold:
            new_data.append(tweet)
    return new_data
```

```
subtweets_data = load_data("../data/other_data/subtweets.json")
```

```
non_subtweets_data = load_data("../data/other_data/non_subtweets.json")
```

Show examples

```
print("Subtweets dataset example:")
print(choice(subtweets_data))
```

Subtweets dataset example:
This little girls to weird for me pure retard

```
print("Non-subtweets dataset example:")  
print(choice(non_subtweets_data))
```

Non-subtweets dataset example:
TESTED: "The Golf Infomercial" Wedge Test

Do golf infomercial wedges really work?

VIEW RESULTS: 📊 📊

Find the length of the smaller dataset

```
smallest_length = len(min([subtweets_data, non_subtweets_data], key=len))
```

Cut both down to be the same length

```
subtweets_data = subtweets_data[:smallest_length]
```

```
non_subtweets_data = non_subtweets_data[:smallest_length]
```

```
print("Smallest dataset length: {}".format(len(subtweets_data)))
```

Smallest dataset length: 7837

Prepare data for training

```
subtweets_data = [(tweet, "subtweet") for tweet in subtweets_data]
```

```
non_subtweets_data = [(tweet, "non-subtweet") for tweet in non_subtweets_data]
```

Combine them

```
training_data = subtweets_data + non_subtweets_data
```

Create custom stop words to include generic usernames, URLs, and hashtags, as well as common English first names

```
names_lower = set([name.lower() for name in open("../data/other_data/first_names.txt").read().split("\n")])
```

```
generic_tokens = {"1", "2", "3"}
```

```
stop_words = text.ENGLISH_STOP_WORDS | names_lower | generic_tokens
```

Build the pipeline

```
sentiment_pipeline = Pipeline([
    ("vectorizer", TfidfVectorizer(tokenizer=tokenizer.tokenize,
                                   ngram_range=(1, 3),
                                   stop_words=stop_words)),
    ("classifier", MultinomialNB()))
])
```

K-Folds splits up and separates out 10 training and test sets from the data, from which the classifier is trained and the confusion matrix and classification reports are updated

```
def confusion_matrices(training_data, num_folds=10):
    text_training_data = np.array([row[0] for row in training_data])
    class_training_data = np.array([row[1] for row in training_data])
    kf = KFold(n_splits=num_folds, random_state=42, shuffle=True)

    cnf_matrix_test = np.zeros((2, 2), dtype=int)
    cnf_matrix_train = np.zeros((2, 2), dtype=int)

    test_reports = []
    train_reports = []
    for i, (train_index, test_index) in enumerate(kf.split(text_training_data)):

        text_train, text_test = text_training_data[train_index], text_training_data[test_index]
        class_train, class_test = class_training_data[train_index], class_training_data[test_index]

        sentiment_pipeline.fit(text_train, class_train)

        predictions_test = sentiment_pipeline.predict(text_test)
        predictions_train = sentiment_pipeline.predict(text_train)

        cnf_matrix_test += confusion_matrix(class_test, predictions_test)
        cnf_matrix_train += confusion_matrix(class_train, predictions_train)

        print("Test Data Iteration {}".format(i+1))

        test_report = classification_report(class_test, predictions_test, digits=4)
        test_reports.append(test_report)
        print(test_report)

        print(("Test Data Null Accuracy: {:.4f}\n"
              .format(max(pd.value_counts(pd.Series(class_test))/float(len(class_test))))))
        print(("Test Data Accuracy: {:.4f}\n"
              .format(accuracy_score(class_test, predictions_test))))
        print("="*53)

        print("Train Data Iteration {}".format(i+1))

        train_report = classification_report(class_train, predictions_train, digits=4)
        train_reports.append(train_report)
        print(train_report)

        print(("Train Data Null Accuracy: {:.4f}\n"
              .format(max(pd.value_counts(pd.Series(class_train))/float(len(class_train))))))
        print(("Train Data Accuracy: {:.4f}\n"
              .format(accuracy_score(class_train, predictions_train))))
        print("="*53)

    def reports_mean(reports):
        reports_lists_of_strings = [report.split("\n") for report in reports]
        reports = [[float(e) for e in report_string[2][16:].split()],
                   [float(e) for e in report_string[3][16:].split()],
                   [float(e) for e in report_string[5][16:].split()]]
                    for report_string in reports_lists_of_strings]
        mean_list = np.mean(np.array(reports), axis=0).tolist()
        print("           precision   recall   f1-score   support")
        print()
        print("non-subtweet      {0:.4f}      {1:.4f}      {2:.4f}      {3:d}"
              .format(mean_list[0][0],
                      mean_list[0][1],
```

```

        mean_list[0][2],
        int(mean_list[0][3]))
    print("    subtweet      {0:.4f}    {1:.4f}    {2:.4f}    {3:d}".format(mean_list[1][0],
        mean_list[1][1],
        mean_list[1][2],
        int(mean_list[1][3])))

    print()
    print(" avg / total      {0:.4f}    {1:.4f}    {2:.4f}    {3:d}".format(mean_list[2][0],
        mean_list[2][1],
        mean_list[2][2],
        int(mean_list[2][3])))

    print()
    print("="*53)

    print("Test Data Averages Across All Folds:")
    reports_mean(test_reports)
    print("Train Data Averages Across All Folds:")
    reports_mean(train_reports)
    return {"Test": cnf_matrix_test, "Train": cnf_matrix_train}

```

```

%%time
cnf_matrices = confusion_matrices(training_data)
cnf_matrix_test = cnf_matrices["Test"]
cnf_matrix_train = cnf_matrices["Train"]

```

```

Test Data Iteration 1:
      precision    recall  f1-score   support

non-subtweet      0.7338      0.6431      0.6855        793
   subtweet      0.6758      0.7613      0.7160        775

 avg / total      0.7052      0.7015      0.7006       1568

```

Test Data Null Accuracy: 0.5057

Test Data Accuracy: 0.7015

```

=====
Train Data Iteration 1:
      precision    recall  f1-score   support

non-subtweet      0.9907      0.9806      0.9856       7044
   subtweet      0.9808      0.9908      0.9858       7062

 avg / total      0.9857      0.9857      0.9857      14106

```

Train Data Null Accuracy: 0.5006

Train Data Accuracy: 0.9857

```

=====
Test Data Iteration 2:
      precision    recall  f1-score   support

non-subtweet      0.6940      0.6324      0.6618        789
   subtweet      0.6584      0.7176      0.6867        779

 avg / total      0.6763      0.6747      0.6742       1568

```

Test Data Null Accuracy: 0.5032

Test Data Accuracy: 0.6747

```

=====
Train Data Iteration 2:
      precision    recall  f1-score   support

non-subtweet      0.9908      0.9786      0.9847       7048

```

| | | | | |
|-------------|--------|--------|--------|-------|
| subtweet | 0.9789 | 0.9909 | 0.9849 | 7058 |
| avg / total | 0.9848 | 0.9848 | 0.9848 | 14106 |

Train Data Null Accuracy: 0.5004

Train Data Accuracy: 0.9848

=====

Test Data Iteration 3:

| | | | | |
|--------------|-----------|--------|----------|---------|
| | precision | recall | f1-score | support |
| non-subtweet | 0.7021 | 0.6866 | 0.6943 | 769 |
| subtweet | 0.7047 | 0.7196 | 0.7121 | 799 |
| avg / total | 0.7034 | 0.7034 | 0.7033 | 1568 |

Test Data Null Accuracy: 0.5096

Test Data Accuracy: 0.7034

=====

Train Data Iteration 3:

| | | | | |
|--------------|-----------|--------|----------|---------|
| | precision | recall | f1-score | support |
| non-subtweet | 0.9869 | 0.9829 | 0.9849 | 7068 |
| subtweet | 0.9829 | 0.9869 | 0.9849 | 7038 |
| avg / total | 0.9849 | 0.9849 | 0.9849 | 14106 |

Train Data Null Accuracy: 0.5011

Train Data Accuracy: 0.9849

=====

Test Data Iteration 4:

| | | | | |
|--------------|-----------|--------|----------|---------|
| | precision | recall | f1-score | support |
| non-subtweet | 0.7313 | 0.6355 | 0.6800 | 801 |
| subtweet | 0.6651 | 0.7562 | 0.7077 | 767 |
| avg / total | 0.6989 | 0.6945 | 0.6936 | 1568 |

Test Data Null Accuracy: 0.5108

Test Data Accuracy: 0.6945

=====

Train Data Iteration 4:

| | | | | |
|--------------|-----------|--------|----------|---------|
| | precision | recall | f1-score | support |
| non-subtweet | 0.9907 | 0.9802 | 0.9854 | 7036 |
| subtweet | 0.9805 | 0.9908 | 0.9856 | 7070 |
| avg / total | 0.9856 | 0.9855 | 0.9855 | 14106 |

Train Data Null Accuracy: 0.5012

Train Data Accuracy: 0.9855

=====

Test Data Iteration 5:

| | | | | |
|--------------|-----------|--------|----------|---------|
| | precision | recall | f1-score | support |
| non-subtweet | 0.7078 | 0.6560 | 0.6809 | 779 |
| subtweet | 0.6828 | 0.7322 | 0.7067 | 788 |
| avg / total | 0.6952 | 0.6943 | 0.6939 | 1567 |

Test Data Null Accuracy: 0.5029

Test Data Accuracy: 0.6943

```
=====
Train Data Iteration 5:
      precision    recall  f1-score   support

non-subtweet      0.9871      0.9829      0.9849        7058
  subtweet      0.9829      0.9871      0.9850        7049

 avg / total      0.9850      0.9850      0.9850       14107
```

Train Data Null Accuracy: 0.5003

Train Data Accuracy: 0.9850

```
=====
Test Data Iteration 6:
      precision    recall  f1-score   support

non-subtweet      0.6836      0.6583      0.6707         758
  subtweet      0.6906      0.7145      0.7023         809

 avg / total      0.6872      0.6873      0.6870       1567
```

Test Data Null Accuracy: 0.5163

Test Data Accuracy: 0.6873

```
=====
Train Data Iteration 6:
      precision    recall  f1-score   support

non-subtweet      0.9874      0.9846      0.9860        7079
  subtweet      0.9845      0.9873      0.9859        7028

 avg / total      0.9860      0.9860      0.9860       14107
```

Train Data Null Accuracy: 0.5018

Train Data Accuracy: 0.9860

```
=====
Test Data Iteration 7:
      precision    recall  f1-score   support

non-subtweet      0.7003      0.6285      0.6625         751
  subtweet      0.6876      0.7525      0.7185         816

 avg / total      0.6937      0.6930      0.6917       1567
```

Test Data Null Accuracy: 0.5207

Test Data Accuracy: 0.6930

```
=====
Train Data Iteration 7:
      precision    recall  f1-score   support

non-subtweet      0.9860      0.9852      0.9856        7086
  subtweet      0.9851      0.9859      0.9855        7021

 avg / total      0.9855      0.9855      0.9855       14107
```

Train Data Null Accuracy: 0.5023

Train Data Accuracy: 0.9855

```
=====
Test Data Iteration 8:
      precision    recall  f1-score   support
```

| | | | | |
|--------------|--------|--------|--------|-----|
| non-subtweet | 0.7342 | 0.6429 | 0.6855 | 812 |
| subtweet | 0.6612 | 0.7497 | 0.7027 | 755 |

| | | | | |
|-------------|--------|--------|--------|------|
| avg / total | 0.6990 | 0.6943 | 0.6938 | 1567 |
|-------------|--------|--------|--------|------|

Test Data Null Accuracy: 0.5182

Test Data Accuracy: 0.6943

=====

Train Data Iteration 8:

| | | | | |
|--|-----------|--------|----------|---------|
| | precision | recall | f1-score | support |
|--|-----------|--------|----------|---------|

| | | | | |
|--------------|--------|--------|--------|------|
| non-subtweet | 0.9906 | 0.9795 | 0.9850 | 7025 |
| subtweet | 0.9799 | 0.9908 | 0.9853 | 7082 |

| | | | | |
|-------------|--------|--------|--------|-------|
| avg / total | 0.9852 | 0.9852 | 0.9852 | 14107 |
|-------------|--------|--------|--------|-------|

Train Data Null Accuracy: 0.5020

Train Data Accuracy: 0.9852

=====

Test Data Iteration 9:

| | | | | |
|--|-----------|--------|----------|---------|
| | precision | recall | f1-score | support |
|--|-----------|--------|----------|---------|

| | | | | |
|--------------|--------|--------|--------|-----|
| non-subtweet | 0.7321 | 0.6429 | 0.6846 | 829 |
| subtweet | 0.6472 | 0.7358 | 0.6886 | 738 |

| | | | | |
|-------------|--------|--------|--------|------|
| avg / total | 0.6921 | 0.6867 | 0.6865 | 1567 |
|-------------|--------|--------|--------|------|

Test Data Null Accuracy: 0.5290

Test Data Accuracy: 0.6867

=====

Train Data Iteration 9:

| | | | | |
|--|-----------|--------|----------|---------|
| | precision | recall | f1-score | support |
|--|-----------|--------|----------|---------|

| | | | | |
|--------------|--------|--------|--------|------|
| non-subtweet | 0.9919 | 0.9796 | 0.9857 | 7008 |
| subtweet | 0.9801 | 0.9921 | 0.9861 | 7099 |

| | | | | |
|-------------|--------|--------|--------|-------|
| avg / total | 0.9860 | 0.9859 | 0.9859 | 14107 |
|-------------|--------|--------|--------|-------|

Train Data Null Accuracy: 0.5032

Train Data Accuracy: 0.9859

=====

Test Data Iteration 10:

| | | | | |
|--|-----------|--------|----------|---------|
| | precision | recall | f1-score | support |
|--|-----------|--------|----------|---------|

| | | | | |
|--------------|--------|--------|--------|-----|
| non-subtweet | 0.7060 | 0.6799 | 0.6927 | 756 |
| subtweet | 0.7116 | 0.7361 | 0.7236 | 811 |

| | | | | |
|-------------|--------|--------|--------|------|
| avg / total | 0.7089 | 0.7090 | 0.7087 | 1567 |
|-------------|--------|--------|--------|------|

Test Data Null Accuracy: 0.5175

Test Data Accuracy: 0.7090

=====

Train Data Iteration 10:

| | | | | |
|--|-----------|--------|----------|---------|
| | precision | recall | f1-score | support |
|--|-----------|--------|----------|---------|

| | | | | |
|--------------|--------|--------|--------|------|
| non-subtweet | 0.9870 | 0.9849 | 0.9859 | 7081 |
| subtweet | 0.9848 | 0.9869 | 0.9859 | 7026 |

| | | | | |
|-------------|--------|--------|--------|-------|
| avg / total | 0.9859 | 0.9859 | 0.9859 | 14107 |
|-------------|--------|--------|--------|-------|

Train Data Null Accuracy: 0.5019

Train Data Accuracy: 0.9859

```
=====
Test Data Averages Across All Folds:
      precision    recall  f1-score   support

non-subtweet      0.7125      0.6506      0.6798        783
   subtweet      0.6785      0.7376      0.7065        783

 avg / total      0.6960      0.6939      0.6933       1567
```

```
=====
Train Data Averages Across All Folds:
      precision    recall  f1-score   support

non-subtweet      0.9889      0.9819      0.9854       7053
   subtweet      0.9820      0.9890      0.9855       7053

 avg / total      0.9855      0.9854      0.9854      14106
```

```
=====
CPU times: user 1min 8s, sys: 1.76 s, total: 1min 10s
Wall time: 1min 12s
```

See the most informative features

[How does "MultinomialNB.coef_" work?](#)

```
def most_informative_features(pipeline, n=10000):
    vectorizer = pipeline.named_steps["vectorizer"]
    classifier = pipeline.named_steps["classifier"]

    class_labels = classifier.classes_

    feature_names = vectorizer.get_feature_names()

    top_n_class_1 = sorted(zip(classifier.coef_[0], feature_names))[:n]
    top_n_class_2 = sorted(zip(classifier.coef_[0], feature_names))[-n:]

    return {class_labels[0]: pd.DataFrame({"Log Probability": [tup[0] for tup in top_n_class_1],
                                           "Feature": [tup[1] for tup in top_n_class_1]}),
            class_labels[1]: pd.DataFrame({"Log Probability": [tup[0] for tup in reversed(top_n_class_2)],
                                           "Feature": [tup[1] for tup in reversed(top_n_class_2)]})}
```

```
%%time
most_informative_features_all = most_informative_features(sentiment_pipeline)
```

```
CPU times: user 1.34 s, sys: 39.3 ms, total: 1.38 s
Wall time: 1.38 s
```

```
most_informative_features_non_subtweet = most_informative_features_all["non-subtweet"]
```

```
most_informative_features_subtweet = most_informative_features_all["subtweet"]
```

```
final_features = most_informative_features_non_subtweet.join(most_informative_features_subtweet,
                                                             lsuffix=" (Non-subtweet)",
                                                             rsuffix=" (Subtweet)")
final_features.to_csv("../data/other_data/most_informative_features.csv")
final_features.head(25)
```

| | Feature (Non-subtweet) | Log Probability (Non-subtweet) | Feature (Subtweet) | Log Probability (Subtweet) |
|----|------------------------|--------------------------------|--------------------|----------------------------|
| 0 | !!& | -12.6618 | . | -7.5300 |
| 1 | !!(| -12.6618 | , | -7.9193 |
| 2 | !!) | -12.6618 | " | -8.0928 |
| 3 | !!. | -12.6618 | people | -8.3903 |
| 4 | !!100 | -12.6618 | ? | -8.4594 |
| 5 | !!15 | -12.6618 | don't | -8.5588 |
| 6 | !!3 | -12.6618 | like | -8.5889 |
| 7 | !!5 | -12.6618 | just | -8.6754 |
| 8 | !!8am | -12.6618 | i'm | -8.6969 |
| 9 | !!:) | -12.6618 | ! | -8.9031 |
| 10 | !!;) | -12.6618 | it's | -8.9727 |
| 11 | !!absolutely | -12.6618 | ... | -9.0431 |
| 12 | !!amazing | -12.6618 | you're | -9.0488 |
| 13 | !!ask | -12.6618 | : | -9.0704 |
| 14 | !!awesome | -12.6618 | know | -9.0928 |
| 15 | !!big | -12.6618 | twitter | -9.1443 |
| 16 | !!bite | -12.6618 | friends | -9.1650 |
| 17 | !!close | -12.6618 | time | -9.2879 |
| 18 | !!collection | -12.6618 | want | -9.2923 |
| 19 | !!come | -12.6618 | u | -9.3004 |
| 20 | !!don't | -12.6618 | really | -9.3518 |
| 21 | !!enter | -12.6618 | shit | -9.3699 |
| 22 | !!epic | -12.6618 | good | -9.4017 |
| 23 | !!extremely | -12.6618 | think | -9.4155 |
| 24 | !!family | -12.6618 | make | -9.4225 |

Define function for visualizing confusion matrices

```
def plot_confusion_matrix(cm, classes=["non-subtweet", "subtweet"],
                          title="Confusion Matrix", cmap=plt.cm.Purples):

    cm_normalized = cm.astype("float") / cm.sum(axis=1)[:, np.newaxis]

    plt.imshow(cm, interpolation="nearest", cmap=cmap)
    plt.colorbar()

    plt.title(title, size=18)

    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=45, fontsize=14)
    plt.yticks(tick_marks, classes, fontsize=14)

    thresh = cm.max() / 2.
```

```

for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
    plt.text(j, i, "{} ({:.0%})".format(cm[i, j], cm_normalized[i, j]),
             horizontalalignment="center", size=16,
             color="white" if cm[i, j] > thresh else "black")

plt.tight_layout()

plt.ylabel("True label", fontsize=14)
plt.xlabel("Predicted Label", fontsize=14)

```

Show the matrices

```



np.set_printoptions(precision=2)

plt.figure(figsize=(6, 6))
plot_confusion_matrix(cnf_matrix_test, title="Test Data Confusion Matrix")

plt.figure(figsize=(6, 6))
plot_confusion_matrix(cnf_matrix_train, title="Train Data Confusion Matrix")

plt.show()

```



Update matplotlib style

```
plt.style.use("fivethirtyeight")
```

Save the classifier for another time

```
joblib.dump(sentiment_pipeline, "../data/other_data/subtweets_classifier.pkl");
```

Print tests for the classifier

```
def process_tweets_for_testing(filenamees):
    dataframes = {}
    for filename in filenamees:
        username = splitext(basename(filename))[0][:-7]
        dataframes[username] = {}

        user_df = pd.read_csv(filename).dropna()
        user_df["Text"] = user_df["Text"].str.replace(hashtags_pattern, "❶")
        user_df["Text"] = user_df["Text"].str.replace(urls_pattern, "❷")
        user_df["Text"] = user_df["Text"].str.replace(at_mentions_pattern, "❸")
        user_df["Text"] = user_df["Text"].str.replace("\u2018", "'")
        user_df["Text"] = user_df["Text"].str.replace("\u2019", "'")
        user_df["Text"] = user_df["Text"].str.replace("\u201c", "\"")
        user_df["Text"] = user_df["Text"].str.replace("\u201d", "\"")
        user_df["Text"] = user_df["Text"].str.replace("&quot;", "\"")
        user_df["Text"] = user_df["Text"].str.replace("&", "&")
        user_df["Text"] = user_df["Text"].str.replace(">", ">")
        user_df["Text"] = user_df["Text"].str.replace("<", "<")

        predictions = sentiment_pipeline.predict_proba(user_df["Text"])[:, 1].tolist()
        user_df["SubtweetProbability"] = predictions

        dataframes[username]["all"] = user_df

        scores = user_df[["SubtweetProbability"]].rename(columns={"SubtweetProbability": username})

        dataframes[username]["scores"] = scores
        dataframes[username]["stats"] = scores.describe()

    return dataframes
```

Load the CSV files

```
filenamees = glob("../data/data_for_testing/friends_data/*.csv")
```

```
%%time
dataframes = process_tweets_for_testing(filenamees)
```

```
CPU times: user 9.09 s, sys: 153 ms, total: 9.24 s
Wall time: 9.52 s
```

Show a random table

```
chosen_username = choice(list(dataframes.keys()))
dataframes[chosen_username]["all"].sort_values(by="SubtweetProbability", ascending=False).head(5)
```

| | Text | Date | Favorites | Retweets | Tweet ID | SubtweetProbability |
|--|------|------|-----------|----------|----------|---------------------|
| | | | | | | |

| | Text | Date | Favorites | Retweets | Tweet ID | SubtweetProbability |
|------|---|---------------------|-----------|----------|--------------------|---------------------|
| 2092 | I hate when people overuse emojis | 2015-06-26 13:01:35 | 0 | 0 | 614478624197091328 | 0.8579 |
| 2137 | Also you don't need to resort to social media 24/7 to complain about your very privileged life ㄣ(ツ)ㄣ | 2015-06-15 17:24:46 | 1 | 0 | 610558590278070272 | 0.8443 |
| 2151 | When I try to be supportive and caring I get ignored and then I'm told I'm not being supportive or caring ㄣ(ツ)ㄣ | 2015-06-13 08:44:07 | 0 | 0 | 609702789896372224 | 0.8366 |
| 2134 | What he doesn't know (unless he stalks my twitter which I know he does) is that I have fake accounts following all his social media | 2015-06-15 17:26:41 | 0 | 0 | 610559074820861953 | 0.8177 |
| 1510 | If you don't have tweet notifications turned on for me are we really friends | 2016-07-14 14:21:21 | 1 | 0 | 753655639465922560 | 0.8076 |

Prepare statistics on tweets

```
tests_df = pd.concat([df_dict["scores"] for df_dict in dataframes.values()], ignore_index=True)
```

```
tests_df.describe()
```

| | adhaardesai | akrapf96 | generatedtext | gothodile | juliaeberry | kayleesue | keithohara |
|-------|-------------|-----------|---------------|-----------|-------------|-----------|------------|
| count | 621.0000 | 2640.0000 | 2066.0000 | 3488.0000 | 4356.0000 | 1939.0000 | 1169.0000 |
| mean | 0.4996 | 0.5086 | 0.5438 | 0.5270 | 0.5187 | 0.4976 | 0.4388 |

| | adhaardesai | akrapf96 | generatedtext | gothodile | juliaeberry | kayleesue | keithohara |
|-----|-------------|----------|---------------|-----------|-------------|-----------|------------|
| std | 0.1059 | 0.1150 | 0.1136 | 0.1086 | 0.1023 | 0.1106 | 0.0981 |
| min | 0.1981 | 0.0953 | 0.1266 | 0.1626 | 0.1522 | 0.0566 | 0.1497 |
| 25% | 0.4291 | 0.4304 | 0.4669 | 0.4538 | 0.4492 | 0.4260 | 0.3733 |
| 50% | 0.4971 | 0.5037 | 0.5417 | 0.5217 | 0.5180 | 0.4981 | 0.4379 |
| 75% | 0.5670 | 0.5847 | 0.6213 | 0.5982 | 0.5843 | 0.5669 | 0.5016 |
| max | 0.8457 | 0.8579 | 0.8497 | 0.8749 | 0.8674 | 0.8766 | 0.8157 |

Plot a histogram with three random users

```

random_choices = sample(list(dataframes.values()), 3)
scores = [df_dict["scores"][df_dict["scores"].columns[0]].tolist()
            for df_dict in random_choices]

fig = plt.figure(figsize=(16, 9))
ax = fig.add_subplot(111)

n, bins, patches = ax.hist(scores,
                            bins="scott",
                            color=["#256EFF", "#46237A", "#3DDC97"],
                            density=True,
                            label=["User 1", "User 2", "User 3"],
                            alpha=0.75)

stats = [df_dict["stats"][df_dict["stats"].columns[0]].tolist()
          for df_dict in random_choices]

line_1 = scipy.stats.norm.pdf(bins, stats[0][1], stats[0][2])
ax.plot(bins, line_1, "--", color="#256EFF", linewidth=2)

line_2 = scipy.stats.norm.pdf(bins, stats[1][1], stats[1][2])
ax.plot(bins, line_2, "--", color="#46237A", linewidth=2)


line_3 = scipy.stats.norm.pdf(bins, stats[2][1], stats[2][2])
ax.plot(bins, line_3, "--", color="#3DDC97", linewidth=2)

ax.set_xticks([float(x/10) for x in range(11)], minor=False)
ax.set_title("Distribution of Subtweet Probabilities In User Accounts", fontsize=18)
ax.set_xlabel("Probability That Tweet is a Subtweet", fontsize=18)
ax.set_ylabel("Percent of Tweets with That Probability", fontsize=18)

ax.legend()

plt.show()

```



Plot a histogram with all of them

First, get some statistics

```
new_tests_df = pd.concat([df_dict["scores"].rename(columns={df_dict["scores"].columns[0]: "SubtweetProbability"})
                          for df_dict in dataframes.values()], ignore_index=True)

new_tests_df_stats = new_tests_df.describe()
```

Then view them

```
new_tests_df_stats
```

| | SubtweetProbability |
|-------|---------------------|
| count | 28632.0000 |
| mean | 0.5133 |
| std | 0.1115 |
| min | 0.0566 |
| 25% | 0.4385 |
| 50% | 0.5093 |
| 75% | 0.5860 |
| max | 0.9091 |

Now plot

```
fig = plt.figure(figsize=(16, 9))
ax = fig.add_subplot(111)

n, bins, patches = ax.hist(new_tests_df["SubtweetProbability"].tolist(),
```

```

bins="scott",
color="#983B59",
edgecolor="black",
density=True,
alpha=0.75)

line = scipy.stats.norm.pdf(bins, new_tests_df_stats["SubtweetProbability"][1],
                             new_tests_df_stats["SubtweetProbability"][2])

ax.plot(bins, line, "--", color="#983B59", linewidth=2)


ax.set_xticks([float(x/10) for x in range(11)], minor=False)
ax.set_title("Distribution of Subtweet Probabilities In All User Accounts", fontsize=18)
ax.set_xlabel("Probability That Tweet is a Subtweet", fontsize=18)
ax.set_ylabel("Percent of Tweets with That Probability", fontsize=18)

ax.legend()

plt.show()

```

No handles with labels found to put in legend.



Statistics on training data

Remove mentions of usernames for these statistics

```

training_data = [(tweet[0]
                  .replace("@", "")
                  .replace("#", "")
                  .replace(" ", "")) for tweet in training_data]

```

Lengths

```

length_data = [len(tweet) for tweet in training_data]

```



```
length_data_for_stats = pd.DataFrame({"Length": length_data, "Tweet": training_data})
```

```
# length_data_for_stats = length_data_for_stats[length_data_for_stats["Length"] <= 280]
```

```
# length_data_for_stats = length_data_for_stats[length_data_for_stats["Length"] >= 5]
```

```
length_data = length_data_for_stats.Length.tolist()
```

Top 5 longest tweets

```
length_data_for_stats.sort_values(by="Length", ascending=False).head()
```

| | Length | Tweet |
|-------|--------|---|
| 8887 | 281 | This Tweet does not endorse the use of Nazi Symbols in any form! I think the image which has been published on social media and MSM is a day or two old. It conjures up strong emotions for many people, My question is simple what meaning do you think is being conveyed by the image? |
| 2198 | 281 | I need to learn how to do this. I ask "how can I help" a lot because I genuinely want to make things better for friends , but this *can* put a burden back upon those who are suffering. Sometimes it may be best to just have exuberant and fearless compassion the same way a pet does |
| 1531 | 281 | hi! I'm not normally v personal like this and I probably won't be at least for a v long time but I thought I'd share this \nwhile I was scrolling on Twitter today I had like a sudden impulse to just dump all my thoughts about what id been reading and seeing and so far it actually- |
| 10533 | 281 | Some people are undecided about testing on animals. Understandable. There's so much propaganda and secrecy about it. Here's a quick test though, & you're answer should tell you. What would you do if some man came to your house & squirted disinfectant in your beautiful dog's eyes? |
| 10521 | 281 | Enthralled by Raja Shiv Chhatrapati, a well mounted magnum opus on life of the Maratha warrior at Red Fort. Vividly brought out his philosophies, struggles, inspiration from mother Jijayee & penchant for gender equality through well conceived music, dance & dialogues. A must see! |

Top 5 shortest tweets

```
length_data_for_stats.sort_values(by="Length", ascending=True).head()
```

| | Length | Tweet |
|------|--------|-------|
| 7699 | 1 | A |
| 3473 | 2 | no |
| 5896 | 2 | uh |
| 6676 | 2 | i- |
| 2038 | 2 | Ha |

Tweet length statistics

```
length_data_for_stats.describe()
```

| | Length |
|-------|------------|
| count | 15674.0000 |
| mean | 106.8089 |
| std | 73.8680 |
| min | 1.0000 |
| 25% | 48.0000 |
| 50% | 87.0000 |
| 75% | 150.0000 |
| max | 281.0000 |

Punctuation

```
punctuation_data = [len(set(punctuation).intersection(set(tweet))) for tweet in training_data]
```

```
punctuation_data_for_stats = pd.DataFrame({"Punctuation": punctuation_data, "Tweet": training_data})
```

Top 5 most punctuated tweets

```
punctuation_data_for_stats.sort_values(by="Punctuation", ascending=False).head()
```

| | Punctuation | Tweet |
|-------|-------------|--|
| 8957 | 11 | Going to go ahead and crown myself the absolute emperor of finding things on menus that sound interesting, deciding I would like to try them, then being told "I'm sorry sir, that's actually not available..." [then why the @#\$% is it ON YOUR MENUUUUUUUUU--] |
| 6725 | 9 | 4-yo: DADDEEEEEEE!? LET'S PLAY! Me: Ok, baby. 4yo: you play w/ her. put a dress on her DADDEEEEEEE. Me: Ok. *puts doll in dollhouse* 4yo: SHE DOESN'T GO THERE!! |
| 11718 | 9 | Self-employed people: have you ever turned to social media to call out a client who is many weeks/months delinquent on a payment? (Obviously, you're probably burning a bridge with that move, but if they don't pay...) |
| 13365 | 9 | Billboard Hot 100: (-3) Tell Me You Love Me, [19 weeks]. *peak: * |
| 11845 | 9 | Tucker Carlson Tonight & TFW you're asking about America but you're scolded it's really about Israel ... Tucker: "What is the American national security interest ... in Syria?" Sen. Wicker(R): "Well, if you care about Israel ..." That was the exact question & answer Shocking |

Tweets punctuation statistics

```
punctuation_data_for_stats.describe()
```

| | Punctuation |
|-------|-------------|
| count | 15674.0000 |
| mean | 1.9168 |

| | Punctuation |
|-----|-------------|
| std | 1.5787 |
| min | 0.0000 |
| 25% | 1.0000 |
| 50% | 2.0000 |
| 75% | 3.0000 |
| max | 11.0000 |

Stop words

```
stop_words_data = [len(set(stopwords.words("english")).intersection(set(tweet.lower()))
                    for tweet in training_data]
```

```
stop_words_data_for_stats = pd.DataFrame({"Stop words": stop_words_data, "Tweet": training_data})
```

Top 5 tweets with most stop words

```
stop_words_data_for_stats.sort_values(by="Stop words", ascending=False).head()
```

| | Stop words | Tweet |
|------|------------|---|
| 0 | 8 | I don't yet have adequate words to do so, but someday I wanna write about the beautiful dance which happens in Google docs between a writer & a good editor working simultaneously towards a deadline. When it's working, it's a beautiful dance—though no one really sees it. |
| 9063 | 8 | Honestly yea i fucked up but all of you are trash asf and your opinions mean nothing to me because mother fucker i can fix shit but yall are to close minded to see. |
| 9035 | 8 | The role of DAG Rod Rosenstein will be an Oscar winner in the future film about the Trump presidency. I'd like the story of the first few months to be told through the eyes of the bewildered Sean Spicer. |
| 9038 | 8 | Done watching 'Hacksaw Ridge'. If there's one thing I learned from that movie, it is simply, Have Faith in God. |
| 9039 | 8 | I feel people who can't celebrate or at the very least respect Cardi B's success have never watched the grind from the ground up. They can't understand that her work ethic has gotten her where she is now. You don't have to stand for what's she's about but she's worked for it |

Top 5 tweets with fewest stop words

```
stop_words_data_for_stats.sort_values(by="Stop words", ascending=True).head()
```

| | Stop words | Tweet |
|-------|------------|----------|
| 3632 | 0 | ... |
| 8290 | 0 | 24 |
| 11925 | 0 | FUCK |
| 10940 | 0 | 78 ... ! |

| | Stop words | Tweet |
|------|------------|--------|
| 1796 | 0 | fuck u |

Tweets stop words statistics

```
stop_words_data_for_stats.describe()
```

| | Stop words |
|-------|------------|
| count | 15674.0000 |
| mean | 7.1515 |
| std | 1.3116 |
| min | 0.0000 |
| 25% | 7.0000 |
| 50% | 8.0000 |
| 75% | 8.0000 |
| max | 8.0000 |

Unique words

```
unique_words_data = [len(set(tokenizer.tokenize(tweet))) for tweet in training_data]
```

```
unique_words_data_for_stats = pd.DataFrame({"Unique words": unique_words_data, "Tweet": training_data})
```

```
# unique_words_data_for_stats = unique_words_data_for_stats[unique_words_data_for_stats["Unique words"] >= 2]
```

```
unique_words_data = unique_words_data_for_stats["Unique words"].tolist()
```

Top 5 tweets with most unique words

```
unique_words_data_for_stats.sort_values(by="Unique words", ascending=False).head()
```

| | Tweet | Unique words |
|-------|---|--------------|
| 13936 | GIVE AWAY!\n\nThe rules are really easy, all you have to do is :\n1. Must be following me (i check) \n2. RT and fav this tweet\n3. tag your mutuals/anyone\n4. only 1 winner! \n5. i ship worldwide;) \n\nit ends in 8th May 2018 or when this tweet hit 2k RT and like!\n\nGood luck! ❤️ | 59 |
| 4881 | got into a tepid back nd forth w/ a uknowwhoAJ+columnist bc i said they steal their "hot takes" from blk twitter & alike. wallahi my bdeshi ass did not sign up 4 this app to be called asinine by a 30yrolld pakistani whos whole politics is Post Colonial Memes for Oriental Minded T- | 57 |
| 7013 | Crazy how wrong u can be about someone. A girl I graduated w/ was always doing drugs& got pregnant at 16. I assumed she'd end up being a loser but it turn out she now has 4 beautiful kids& is making over \$4,500/month just off of child support payments from the 3 different dads | 57 |

| | Tweet | Unique words |
|-------|---|--------------|
| 4992 | Got into an argument w/ someone I went to HS w/ & I would js like to repeat again tht I cannot wait to stunt on all the ppl who were bitches to me in HS @ our reunion. Catch me rollin up w/ my sexy ass gf, a nice car, a bomb body & the career of my dreams as a big fuck u to them | 55 |
| 11542 | Thought I'd bring this back... and no, I'm not talking about myself here. I wish just once I'd be so bored with my life that I'd find the time to bash people/celebs I don't like.. I mean if I despise someone THAT much, why still watch his/her every move? 🙄 | 55 |

Top 5 tweets with fewest unique words

```
unique_words_data_for_stats.sort_values(by="Unique words", ascending=True).head()
```

| | Tweet | Unique words |
|-------|-------------------------|--------------|
| 6106 | Annoying | 1 |
| 2525 | Bitch | 1 |
| 12087 | Chandler | 1 |
| 14559 | Yes yes yes yes yes yes | 1 |
| 14442 | Hello\n | 1 |

Tweets unique words statistics

```
unique_words_data_for_stats.describe()
```

| | Unique words |
|-------|--------------|
| count | 15674.0000 |
| mean | 19.2412 |
| std | 11.9298 |
| min | 1.0000 |
| 25% | 10.0000 |
| 50% | 17.0000 |
| 75% | 27.0000 |
| max | 59.0000 |

Plot them

```
length_mean = length_data_for_stats.describe().Length[1]
length_std = length_data_for_stats.describe().Length[2]

fig = plt.figure(figsize=(16, 9))
ax = fig.add_subplot(111)

n, bins, patches = ax.hist(length_data,
                             bins="scott",
                             edgecolor="black",
                             # density=True,
                             color="#12355b",
```

```


        alpha=0.5)

# length_line = scipy.stats.norm.pdf(bins, length_mean, length_std)
# ax.plot(bins, length_line, "--", linewidth=3, color="#415d7b")

ax.set_title("Training Dataset Distribution of Tweet Lengths", fontsize=18)
ax.set_xlabel("Tweet Length", fontsize=18);
ax.set_ylabel("Number of Tweets with That Length", fontsize=18);

plt.show()

```



```

punctuation_mean = punctuation_data_for_stats.describe().Punctuation[1]
punctuation_std = punctuation_data_for_stats.describe().Punctuation[2]

fig = plt.figure(figsize=(16, 9))
ax = fig.add_subplot(111)


n, bins, patches = ax.hist(punctuation_data,
                            bins="scott",
                            edgecolor="black",
                            # density=True,
                            color="#420039",
                            alpha=0.5)

# punctuation_line = scipy.stats.norm.pdf(bins, punctuation_mean, punctuation_std)
# ax.plot(bins, punctuation_line, "--", linewidth=3, color="#673260")

ax.set_title("Training Dataset Distribution of Punctuation", fontsize=18)
ax.set_xlabel("Punctuating Characters in Tweet", fontsize=18)
ax.set_ylabel("Number of Tweets with That Number of Punctuating Characters", fontsize=18)

plt.show()

```



```
stop_words_mean = stop_words_data_for_stats.describe()["Stop words"][1]
stop_words_std = stop_words_data_for_stats.describe()["Stop words"][2]

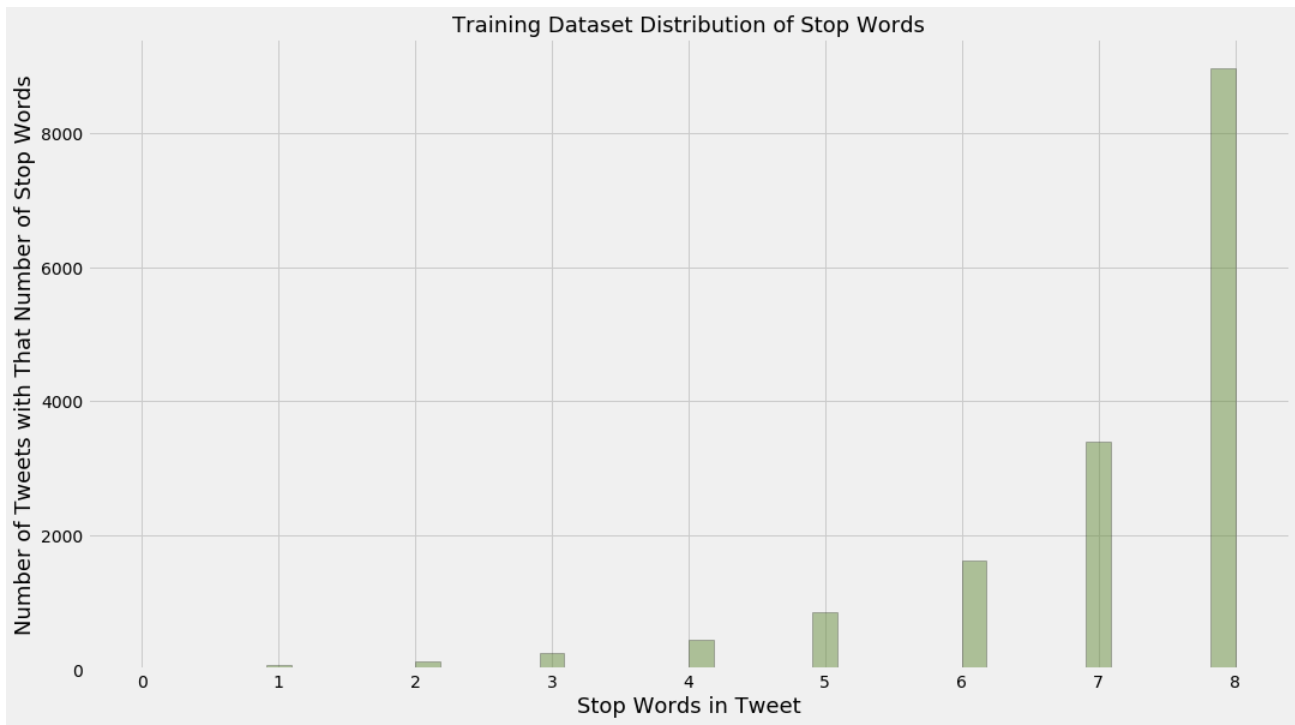
fig = plt.figure(figsize=(16, 9))
ax = fig.add_subplot(111)

n, bins, patches = ax.hist(stop_words_data,
                             bins="scott",
                             edgecolor="black",
                             # density=True,
                             color="#698f3f",
                             alpha=0.5)

# stop_words_line = scipy.stats.norm.pdf(bins, stop_words_mean, stop_words_std)
# ax.plot(bins, stop_words_line, "--", linewidth=3, color="#87a565")

ax.set_title("Training Dataset Distribution of Stop Words", fontsize=18)
ax.set_xlabel("Stop Words in Tweet", fontsize=18)
ax.set_ylabel("Number of Tweets with That Number of Stop Words", fontsize=18)

plt.show()
```



```
unique_words_mean = unique_words_data_for_stats.describe()["Unique words"][1]
unique_words_std = unique_words_data_for_stats.describe()["Unique words"][2]

fig = plt.figure(figsize=(16, 9))
ax = fig.add_subplot(111)

n, bins, patches = ax.hist(unique_words_data,
                             bins="scott",
                             edgecolor="black",
                             # density=True,
                             color="#ca2e55",
                             alpha=0.5)

# unique_words_line = scipy.stats.norm.pdf(bins, unique_words_mean, unique_words_std)
# ax.plot(bins, unique_words_line, "--", linewidth=3, color="#d45776")

ax.set_title("Training Dataset Distribution of Unique Words", fontsize=18)
ax.set_xlabel("Unique Words in Tweet", fontsize=18)
ax.set_ylabel("Number of Tweets with That Number of Unique Words", fontsize=18)

plt.show()
```