

# Don't Take This Personally: Sentiment Analysis for Identification of “Subtweeting” on Twitter

A Senior Project submitted to  
The Division of Science, Mathematics, and Computing  
of  
Bard College

by  
Noah Segal-Gould

Annandale-on-Hudson, New York  
May, 2018



# Abstract

The purpose of this project is to identify subtweets. The Oxford English Dictionary defines “subtweet” as a “[Twitter post] that refers to a particular user without directly mentioning them, typically as a form of furtive mockery or criticism.” This paper details a process for gathering a labeled ground truth dataset, training a classifier, and creating a Twitter bot which interacts with subtweets in real time. The Naive Bayes classifier trained in this project classifies tweets as subtweets and non-subtweets with an average  $F_1$  score of 72%.



# Contents

<b>Abstract</b>	<b>iii</b>
<b>Dedication</b>	<b>vii</b>
<b>Acknowledgments</b>	<b>ix</b>
<b>1 Introduction</b>	<b>1</b>
<b>2 Methods</b>	<b>5</b>
2.1 Gathering Tweets for the Ground Truth Dataset . . . . .	5
2.1.1 The Twitter API . . . . .	5
2.1.2 Searching for Tweets . . . . .	6
2.1.3 Changes in Data Acquisition . . . . .	8
2.2 Language Analysis & Naive Bayes . . . . .	9
2.2.1 Resources for Language Analysis . . . . .	9
2.2.2 Cleaning and Preparing the Data . . . . .	10
2.2.3 TF & TF-IDF . . . . .	12
2.2.4 Naive Bayes . . . . .	13
2.2.5 Statistical Considerations . . . . .	15
2.2.6 K-Folds Cross-Validation . . . . .	16
2.2.7 Training & Testing Naive Bayes . . . . .	17
2.3 Creating the Twitter Bot . . . . .	18
<b>3 Results</b>	<b>21</b>
3.1 Ground Truth Dataset . . . . .	21
3.1.1 Characteristics of the Ground Truth Dataset . . . . .	23
3.2 Statistics . . . . .	26
3.2.1 The Confusion Matrix . . . . .	26
3.2.2 Precision, Recall, & F1 . . . . .	27
3.3 Known Subtweeters . . . . .	27

3.4	Most Informative Features . . . . .	28
3.5	The Twitter Bot . . . . .	29
3.6	Discussion . . . . .	30
<b>4</b>	<b>Conclusion</b>	<b>33</b>
4.1	Summary of Project Achievements . . . . .	33
4.2	Future Work & Considerations . . . . .	33
	<b>Bibliography</b>	<b>35</b>
<b>A</b>	<b>Acquiring Ground Truth Data</b>	<b>37</b>
A.1	subtweets_downloader.py . . . . .	37
A.2	non_subtweets_downloader.py . . . . .	38
<b>B</b>	<b>Training &amp; Testing the Classifier</b>	<b>41</b>
B.1	classifier_creator.py . . . . .	41
B.2	live_subtweets_classifier.py . . . . .	55

# Dedication

I dedicate this senior project to @jack, who has willfully made numerous changes to Twitter which inevitably angered millions.





# Acknowledgments

Thank you professors Sven Anderson, Keith O'Hara, and Rebecca Thomas for making this project possible through your combined efforts to teach and advise me. Thank you Benjamin Sernau '17 for enduring through three years of Computer Science courses with me and being a source of unending joy in my life. Thank you to Julia Berry '18, Aaron Krapf '18, and Zoe Terhune '18 for being my very best friends and giving me things worth caring about. Finally, thank you to my parents Tammy Segal and Emily Taylor for your constant support and patience throughout my four years at Bard College.



# 1

## Introduction

“Subtweet” was coined in December of 2009 by Twitter user Chelsea Rae (Rae, 2009) and was entered into Urban Dictionary the following August (Urban Dictionary, 2010). In “To tweet or subtweet?: Impacts of social networking post directness and valence on interpersonal impressions” (Edwards and Harris, 2016), Edwards and Harris sought to analyze student participants’ perceptions of known subtweeters. In the news, too, subtweets have garnered attention in *The Atlantic* (Madrigal, 2014), *The Washington Post* (Dewey, 2016), and *Slate* (Hassler, 2016). In news media, subtweets garner attention for their prevalence among government officials as well. Following President Donald Trump’s inauguration, The Washington Post compiled its “A running list of all the possible subtweets of President Trump from government Twitter accounts,” (Ohlheiser, 2017) cementing subtweets as particularly newsworthy.

There were over 140 million active Twitter users who sent 340 million text-based tweets to the platform every day by March of 2012 (Twitter, 2012). Since Twitter-founder Jack Dorsey sent the first Tweet in March of 2006 (Dorsey, 2006) social scientists, political scientists, and computer scientists have applied machine learning techniques to understand the patterns and structures of the conversations held on the platform. Through sentiment analysis, we are able to use machine learning to identify patterns within natural language which indicate particular feelings both

broadly (e.g. positive, neutral, or negative) and toward topics (e.g. politics, terrorism, brands, etc.).

On Twitter, the most common way to publicly communicate with another user is to compose a tweet and place an “@” before the username of that user somewhere in the tweet (e.g. ”How are you doing, @NoahSegalGould?”). Through this method, public discussions on Twitter maintain a kind of accountability: even if one were to miss the notification that they were mentioned in a tweet, one’s own dashboard keeps a running list of their most recent mentions.

If an individual sought to disparage or mock another, they could certainly do so directly. But the targeted user would probably notice, and through the search functions of the platform, anyone could see who has mentioned either their own or another’s username. Instead, a phenomenon persists in which users of the platform deliberately insult others in a vague manner by making complaints while omitting the targets of those complaints.

Although the OED’s definition states that a subtweet “...refers to a particular user without directly mentioning them, typically as a form of furtive mockery or criticism,” it is perhaps too restrictive. Some individuals believe subtweets abide by this definition, but others expand it to allow inclusion of others’ real names (especially if that individual does not own a Twitter account), and some do not even require that a particular user be the target of the tweet. Because subtweeting is colloquial in nature, we will expand the definition of subtweet to permit these less restrictive features.

Sentiment analysis on social networking services such as Twitter has garnered attention within seemingly distinct fields of interest. In “Text mining for market prediction: A systematic review,” Nassirtoussi et al. surveyed varied methods for text-mining social media for sentiment analysis of financial markets and approached that problem with both behavioral and economic considerations in mind (Nassirtoussi et al., 2014). Following a terrorist event in Woolwich, London in 2013, Burnap et al. analyzed the immediate Twitter response following the attack to inform statistics on how long it takes for responses from official sources to disseminate during crises (Burnap et

al., 2014). Prior research of these kinds utilizes sentiment analysis techniques on tweets, but no known research exists which specifically performs any sentiment analysis on subtweets.

Long before Twitter, psychologist Gordon Allport wrote about “antilocution” in *The Nature of Prejudice* (Allport, 1954). For Allport, antilocution was the first of several degrees of apathy which measure prejudice in a society. It represented the kind of remarks which target a person, group, or community in a public or private setting but do not address the targeted individual directly. Different from both hate speech and subtweeting, antilocution necessitates that an in-group ostracize an unaware out-group through its biases.

The most germane research available focuses on sentiment analysis of figurative language. Determining sentiment based on features of text which are distinctly separate from their literal interpretations presents difficulties for human readers as well as computer programs. In *SemEval*, the International Workshop on Semantic Evaluation, analysis of figurative language on Twitter has been a core task for their competition since 2015 (Ghosh et al., 2015) and returns this year with a specific focus on ironic tweets (Van Hee et al., 2018). In this year’s description for “Task 3: Irony detection in English tweets,” Van Hee et al. touch upon online harassment as a potential point of significance for sentiment analysis of ironic tweets.

We pursue sentiment analysis of subtweets in order to challenge the hypocrisy of utilizing a service which presents itself as a public forum to speak in distinctly private ways. Toward this end, these are our goals: this project will provide a framework for collecting examples of subtweets, train a classification algorithm using those examples, and finally utilize that classifier in real time to make tweets which were intended to be unseen by specific parties easily accessible to all parties. In presenting covertly hurtful content as obviously hurtful in a public fashion, perhaps it will promote a particular awareness that tweets posted by public accounts are indeed publicly accessible, and that Twitter’s Terms of Service (Twitter, 2016) allows for this kind of monitoring.

Twitter bots are automated programs which access and archive Twitter. During the 2016 presidential election, The Atlantic’s Oxford University researchers revealed that more than 33% of

pro-Trump tweets and nearly 20% of pro-Clinton tweets came from automated Twitter accounts between the first and second presidential debates (Guilbeault and Woolley, 2016). Sometimes indistinguishable to untrained users, Twitter bots often misuse the API and violate the Terms of Service to harass or serve an unseen agenda. In contrast, some Twitter bots are purely entertaining like “Pentametron,” a Twitter bot that searches Twitter every few hours for tweets which happen to fit the rhythm of iambic pentameter (Bhatnagar, 2018).

Using a machine-learning approach to perform sentiment analysis, syntactic and linguistic features are typically utilized in probabilistic (e.g. Naive Bayes and Maximum Entropy) and linear (e.g. Support Vector Machines and Neural Networks) classification algorithms. The probabilistic approach is sometimes called *generative* because such models generate the probabilities of sampling particular terms (Medhat et al., 2014). Linear classification utilizes the vectorized feature space of words, sentences, or documents to find a separating hyperplane between multiple classes. In this project, we approach the problem of identifying subtweets using the probabilistic Naive Bayes classification algorithm and create a Twitter bot to interact with identified subtweets.

# 2

## Methods

### 2.1 Gathering Tweets for the Ground Truth Dataset

#### *2.1.1 The Twitter API*

Twitter provides a free Application Programming Interface (API) to registered users and has done so since September of 2006 (Stone, 2006). The API allows developers to programmatically access and influence tweets individually or through real time search filters, and also read and write direct messages (Twitter, 2018). The creation of a Twitter application which utilizes the API requires creation and email verification of an account, and developers are also required to agree to the terms of service (Twitter, 2016). Creation of an application provides developers with authentication tokens which can then be used to access the API.

To make creation of Twitter applications easier, Tweepy (Roesslein, 2009) is an open source library for the Python programming language which provides methods and classes used to interact with the API and its status objects (Twitter, 2018). A Twitter status object is a dictionary of key and value pairs which contains text, media, and user information associated with particular tweets. There are rate limits for both reading and writing to the API which must be kept in mind when programming for it.

### 2.1.2 Searching for Tweets

For acquisition of a ground truth dataset, we consider **true subtweets** to be tweets to which another Twitter user replied who specifically called it out as a subtweet. We consider **true non-subtweets** to be tweets to which another user replied who specifically did **not** call it out as a subtweet.

```

1 def get_subtweets(max_tweets=2500,
2                   query=("subtweet AND @ since:2018-04-01"
3                         "exclude:retweets filter:replies")):
4     subtweets_ids_list = []
5     subtweets_list = []
6     for potential_subtweet_reply in tweepy.Cursor(api.search, lang="en",
7                                                   tweet_mode="extended",
8                                                   q=query).items(max_tweets):
9         potential_subtweet_original = first_tweet(potential_subtweet_reply)
10        if (not potential_subtweet_original.in_reply_to_status_id_str
11            and potential_subtweet_original.user.lang == "en"):
12            if (potential_subtweet_original.id_str in subtweets_ids_list
13                or "subtweet" in potential_subtweet_original.full_text
14                or "Subtweet" in potential_subtweet_original.full_text
15                or "SUBTWEET" in potential_subtweet_original.full_text):
16                continue
17            else:
18                subtweets_ids_list.append(potential_subtweet_original.id_str)
19                subtweets_list.append({"tweet_data": potential_subtweet_original._json,
20                                     "reply": potential_subtweet_reply._json})
21            with open("../data/other_data/subtweets.json", "w") as outfile:
22                json.dump(subtweets_list, outfile, indent=4)
23    return subtweets_list

```

After the API credentials are loaded, this Python code is used to download tweets with replies which **do** and **do not** call them out as subtweets. Tweepy provides the `Cursor` object which is used to iterate through different sections of the Twitter API. In this example, we use it with Twitter’s search API to find tweets matching our query. This particular version finds **true subtweets** with their accusatory replies, but it requires only modification of its query argument to download **true non-subtweets**. Within the `for` loop, we confirm that the following qualities hold true:

- The alleged subtweet or non-subtweet is not in reply to any other tweet.
- The user who posted it used English as their primary language.
- We do not download duplicate tweets.
- It does not contain the string “subtweet” (with variations in capitalization).



For acquisition of both subtweets and non-subtweets, this function can be modified to accept different queries. Respectively, these are the queries which were utilized in acquiring so-called subtweet-accusations (i.e. these replies typically claim that the tweet to which they reply was a subtweet) and non-subtweet-accusations (i.e. these replies are essentially normal replies to normal tweets):

```
"subtweet AND @ since:2018-04-01 exclude:retweets filter:replies"
```

```
"-subtweet AND @ since:2018-04-01 exclude:retweets filter:replies"
```

The only difference between the two is that the former searches for tweets which contain the string “subtweet” and the latter searches for tweets which exclude it. Both access the API as far back as it will allow (typically one week for free API users) and exclude retweets (i.e. unoriginal tweets which are reposted) while specifically searching for tweets which were in reply to other tweets. Unfortunately, tweet status objects representing replies to tweets do not contain the object data of the tweet to which they reply.

```
1 def first_tweet(tweet_status_object):
2     try:
3         return first_tweet(api.get_status(tweet_status_object.in_reply_to_status_id_str,
4                                           tweet_mode="extended"))
5     except tweepy.TweepError:
6         return tweet_status_object
```

To obtain this object, this code essentially goes up a chain of recursively finding the original tweet to which a reply was made until the API can no longer find another tweet object with an `in_reply_to_status_id_str` attribute, indicating that the tweet is an original **true subtweet** or **true non-subtweet**. The two versions of this program for acquiring **true subtweets** and **true non-subtweets** ran for three weeks between March and April of 2018, acquiring over 20,000 tweets which compose our ground truth dataset.

	True Subtweet Data	True Non-Subtweet Data
Tweet	Talk to him again about “dropping me” and you’ll get your teeth knocked out	That’s been one of my biggest issues here; the onus is on ordinary people who, in their spare time, must campaign for the basic services of a city. This is not how progressive cities should be built
Reply	Thomas don’t subtweet me during work hours	i guess i am not as creative as i thought

This table features a **true subtweet** and its associated reply and a **true non-subtweet** and its associated reply.

### 2.1.3 *Changes in Data Acquisition*

This approach for creating a ground truth dataset relies on a particular phenomenon in which Twitter users call-out the subtweets of their peers. The following pattern was observed: a user posts a subtweet which is easily recognized by a peer, and that peer then replies to that tweet in order to complain that the original user was subtweeting or to ask if the tweet was indeed a subtweet. Initially, the program for acquiring the ground truth dataset used the Twitter API’s search functionality to specifically search for replies to tweets which contained some form of the string “subtweet.” It utilized the API’s status object to access the tweet to which it was replying. For 73 days, each day’s alleged subtweets and their associated accusatory replies were saved.

Previously, the classifier was trained using a dataset which was half composed of these alleged subtweets and half composed of tweets randomly selected from a pre-labeled sentiment analyzed tweets dataset (Go et al., 2009). This procedure failed to make the training data representative of **true subtweets** and **true non-subtweets**. The alleged subtweets downloading program was revised and it has been set to download tweets with replies which specifically do **not** contain the string “subtweet.” In both the program which downloads subtweets and the program which downloads non-subtweets, the assumptions about these interactions will not hold true in every case. They are intended as generalizations which make acquiring a ground truth dataset for use in performing binary classification significantly easier and less time-consuming than finding and labeling subtweets and non-subtweets by hand. Indeed, the dataset utilized by Go et al. uses a similar method for acquiring labeled data. In their *Sentiment140* dataset, the labels were acquired according to emoticons present within the tweets instead of through hand-labeling by actual humans.

## 2.2 Language Analysis & Naive Bayes

Before training the classifier, the ground truth dataset is modified such that the important features within it are easily accessible. Changes are made to preserve the characteristics of the text which are relevant to the goal of the analysis and to leave out the ones which are irrelevant. Because we will be using Naive Bayes, we must keep in mind which features in each tweet (e.g. URLs and user names) ought to influence the probabilities that an entire feature-set (i.e. that whole tweet) suits a particular class.

### *2.2.1 Resources for Language Analysis*

#### Regular Expressions

For text classification through machine learning, it is popular to modify the ground truth dataset to make features which are not important to the classification problem as **generic** as possible. For classification of subtweets, the classifier will treat URLs, mentions of usernames, and English first names generically. In other words, it will keep track of the existence of those features but specifically will not encounter the text contained within them. In identification of subtweets, there exists no syntactic or linguistic significance in the format of a URL or the name a user chooses to associate with themselves or another. However, the existence of those features within the tweet remains important. For this kind of substring searching, pattern matching through regular expressions was used to replace every occurrence of URLs, usernames, and first names with special tokens which were not already in the dataset. The top 100 most common English names for both men and women over the last century were acquired from the United States Department of Social Security.

#### Tokenization

Instead of training the classifier on entire strings, **tokenization** is necessary in order to extract individual features from the text. The Natural Language Toolkit (Bird and Loper, 2004) provides a tweet tokenizer to achieve this. For some string, the tokenizer splits apart words, usernames, URLs, hashtags, and punctuating characters as individual tokens. NLTK's tweet tokenizer also

appropriately distinguishes between punctuating characters and emoticons composed of punctuating characters.

### N-Grams

An **n-gram** is a contiguous sequence of  $n$  tokens in a piece of text. For example, given a string such as “This is a test,” the bigrams ( $n = 2$ ) for this string are “This is,” “is a,” and “a test.” Instead of training the classifier using unigrams ( $n = 1$ ) exclusively, we train it using unigrams, bigrams, and trigrams ( $n = 3$ ). Thus, when the probability that some specific token within a tweet belongs to a specific class is calculated, its neighbors are also considered in combination with it. N-grams enable the classifier to treat particular groupings of tokens with some size  $n$  as importantly as it treats the individual tokens, thus identifying particular word groupings most associated with the classes.

### Stop Words

A list of stop words typically contains the most common words in a language. For English text, the list is often composed of words such as “the,” “it,” and “of.” Tokens matching stop words are ignored during classifier training because they are too common to help the classifier distinguish subtweets from non-subtweets.

#### 2.2.2 Cleaning and Preparing the Data

Although **true subtweets** and **true non-subtweets** are identified using characteristics of the Twitter API’s tweet status objects, we utilize Naive Bayes exclusively for text classification on the unicode text contained within each alleged subtweet and non-subtweet status object. Thus, we ignore the replies to these tweets.

```

1 def load_data(filename, threshold=0.1):
2     data = [(urls_pattern.sub("GENERIC_URL",
3         at_mentions_pattern.sub("GENERIC_MENTION",
4         names_pattern.sub("GENERIC_NAME",
5         t["tweet_data"]["full_text"])))
6         .replace("\u2018", "'')
7         .replace("\u2019", "'')
8         .replace("\u201c", "\"")
9         .replace("\u201d", "\"")
10        .replace("&quot;", "\"")
11        .replace("&", "&")
12        .replace(">", ">")
13        .replace("<", "<")

```

```

14         for t in json.load(open(filename))
15         if t["tweet_data"]["lang"] == "en"
16         and t["reply"]["lang"] == "en"
17         and t["tweet_data"]["user"]["lang"] == "en"
18         and t["reply"]["user"]["lang"] == "en"]
19     new_data = []
20     for tweet in data:
21         tokens = tokenizer.tokenize(tweet)
22         english_tokens = [english_dict.check(token) for token in tokens]
23         percent_english_words = sum(english_tokens)/len(english_tokens)
24         if percent_english_words >= threshold:
25             new_data.append(tweet)
26     return new_data

```

The `load_data` function genericizes URLs, mentions of user names, and mentions of English first names. Using regular expressions for pattern matching, these substrings are replaced with special identifiers. The tweets are also cleaned to make HTML characters and unicode characters more consistent. The list comprehension intentionally excludes **non-English** tweets and those which were **not** posted by accounts which list their primary language as English. NLTK's tweet tokenizer is utilized at the end to check for tweets which contain at least 10% English tokens. This language detection is performed on each token in the tweet using the pyEnchant library (Kelly, 2016) which primarily serves as a spell-checker. The resulting dataset of either **true subtweets** or **true non-subtweets** is returned as a list.

After text cleaning, we remove tweets which are present in both the dataset of **true subtweets** and **true non-subtweets**. Duplicates may have appeared because one user thought a tweet was a subtweet but another did not.

```

1     subtweets_data = load_data("../data/other_data/subtweets.json")
2     non_subtweets_data = load_data("../data/other_data/non_subtweets.json")
3     subtweets_data = [tweet for tweet in subtweets_data
4                       if tweet not in non_subtweets_data]
5     non_subtweets_data = [tweet for tweet in non_subtweets_data
6                          if tweet not in subtweets_data]

```

With the duplicates gone, we limit the size of the larger dataset to be the same as the smaller of the two. Thus, both the **true subtweets** and **true non-subtweets** compose the entire final ground truth dataset equally.

```

1     smallest_length = len(min([subtweets_data, non_subtweets_data], key=len))
2
3     subtweets_data = sample(subtweets_data, smallest_length)
4     non_subtweets_data = sample(non_subtweets_data, smallest_length)
5
6     subtweets_data = [(tweet, "subtweet") for tweet in subtweets_data]

```

```

7         non_subtweets_data = [(tweet, "non-subtweet") for tweet in non_subtweets_data]
8
9         training_data = subtweets_data + non_subtweets_data

```

In this code, Python’s `random` library provides the `sample` function which randomly returns a list of  $n$  items from a list. After both datasets are made the same length, the lists of strings are made into lists of tuples. For each tuple, the **true subtweet** or **true non-subtweet** is associated with its label (i.e. subtweet and non-subtweet). Finally, these two lists of tuples are put together for use in training the classifier.

### 2.2.3 TF & TF-IDF

Term frequency (TF) is a simple method for vectorizing text in which all terms (i.e. tokens, features, words, etc.) in the corpus are featured in a vector for each document, and the frequency of each term is reflected in the number representing the corresponding term. The bag of words model essentially vectorizes features using this method. Thus,  $tf(t, d) = f_{t,d}$ .

Unfortunately, TF falls short when the corpus of documents contains terms which appear frequently but do not necessarily help inform the classifier on terms that are best associated with a particular class. TF-IDF, or term frequency-inverse document frequency, is the product of the TF for a specific term and the inverse document frequency (IDF) for that same term. The TF is equal to the ratio between the number of occurrences of a term in a document, and the total number of words in that document. IDF, then, is the logarithm of the ratio between the number of documents in the corpus, and the number of documents which contain that term. The product of TF and IDF assigns weights which appropriately value terms which are frequent within a document but rare in the entire corpus of documents. Thus,  $idf = \log \frac{N}{n_t}$  where  $N$  is the total number of documents and  $n_t$  is the number of documents which contain the term. Because the weighted feature vectors calculated using TF-IDF follow a multinomial distribution, our classification algorithm is specifically a Multinomial Naive Bayes classifier.

### 2.2.4 Naive Bayes

Naive Bayes stands out as particularly simple and common for use in text classification. A **bag of words** model typically ignores word positions in favor of keeping track of raw token frequencies, which are weighted to produce TF-IDF feature vectors. Then, Bayes theorem is utilized to predict the probability that a given feature set (e.g. words, sentences, etc.) belongs to a particular label (i.e. a category or class). Bayes theorem is a means of predicting the **posterior probability** that an event occurs given the observation of another event. It relies on the **conditional probability** that the other event occurs given the observation of the event and the **prior probabilities** that both events occur in general.

Thus,

$$P(A|B) = \frac{P(B|A)P(A)}{P(B)}$$

where  $P(A|B)$  is the **posterior probability**,  $P(B|A)$  is the **conditional probability**, and  $P(A)$  and  $P(B)$  are the **prior probabilities**.

The **prior probability**  $P(A)$  is otherwise known as the **class probability** and is equal to the general probability of encountering a particular class. In our case, we have chosen to keep our classes balanced (i.e. there are equal numbers of documents in both), so the **class probability** will always be 50%. The other **prior probability**  $P(B)$  or **evidence**, then, refers to the probability of encountering the features within a document **independent** of the class label. The **conditional probability**  $P(B|A)$  refers to the likelihood of encountering the features within a document given that those features belong to a particular class. The **evidence** is often ignored in the final classification step on the basis that it acts merely as a scaling factor when trying to maximize which class produces the greater **posterior probability**. The **posterior probability** is the probability that a particular document belongs to a class given the observed features within that document. The Naive Bayes algorithm maximizes this probability in order to predict which class best fits a document. In all cases where we must calculate the probability for an entire feature-set, we simply take the product of all the extracted feature probabilities in

the document. Consider this classification example:

if  $P(\omega = \text{subtweet} \mid \mathbf{x}) \geq P(\omega = \text{non-subtweet} \mid \mathbf{x})$  classify as subtweet,  
 else classify as non-subtweet.

Dropping the evidence term on the basis that it is constant for both classes, we expand that:

$$P(\omega = \text{subtweet} \mid \mathbf{x}) = P(\mathbf{x} \mid \omega = \text{subtweet}) \cdot P(\text{subtweet})$$

$$P(\omega = \text{non-subtweet} \mid \mathbf{x}) = P(\mathbf{x} \mid \omega = \text{non-subtweet}) \cdot P(\text{non-subtweet})$$

Assuming the **posterior probability** for the former is greater than or equal to the latter, the classifier predicts that the document fits that class. The naive assumption maintains that all features are treated as conditionally independent (i.e. that the presence or omission of a particular feature does not change the likelihood of encountering other features), and although this is frequently violated, Naive Bayes often performs well anyway (Zhang, 2004).

For cases in which the classifier encounters a feature absent from the features which were used to train it, a so-called **zero probability** appears. Because the probability of encountering the feature is 0, **additive smoothing** is often utilized to appropriately weight new features using an extra term  $\alpha$ , so the probability that an entire feature-set fits into a specific class is not 0. In this project, we use **Laplace smoothing** ( $\alpha = 1$ ). This technique smooths categorical data by including the pseudo-count  $\alpha$  into each probability estimate. Thus, the probability of a feature given a particular class becomes:

$$P(x_i \mid \omega_j) = \frac{N_{x_i, \omega_j} + \alpha}{N_{\omega_j} + \alpha d} \quad (i = (1, \dots, d))$$

where  $N_{x_i, \omega_j}$  is the number of times feature  $x_i$  appears in samples from class  $\omega_j$ ,  $N_{\omega_j}$  is the total count of all features in class  $\omega_j$ ,  $\alpha$  is the parameter for additive smoothing, and  $d$  is the dimensionality of the feature vector  $\mathbf{x} = [x_1, \dots, x_d]$ . Compared to datasets which contain millions of tweets such as *Sentiment140* (Go et al., 2009), our classifier has access to significantly fewer documents. We utilize Laplace smoothing because when the classifier is tested on new tweets it will likely encounter never before seen features given the limited size of the ground truth dataset.



### 2.2.5 Statistical Considerations

In the binary classification of subtweets and non-subtweets, we consider true positives (TP) to be **true subtweets** which were correctly labeled as **predicted subtweets**, false positives (FP) to be **true non-subtweets** which were incorrectly labeled as **predicted subtweets**, true negatives (TN) to be **true non-subtweets** which were correctly labeled as **predicted non-subtweets**, and false negatives (FN) to be **true subtweets** which were incorrectly labeled as **predicted non-subtweets**. As such, there are two ways for the classifier to be wrong: it can produce false negatives and false positives. Using TP, FP, TN, and FN, we can calculate statistical measurements on the performance of our classifier such as **precision**, **recall**,  $F_1$  **score**, and **null accuracy**.

Precision

**Precision** refers to the ratio between the true positives, and the true positives and false positives. It is also known as the **positive predictive value**.

$$P = \frac{TP}{TP + FP}$$

Recall

**Recall**, then, refers to the ratio between the number of true positives, and the true positives and false negatives. It is also known as the **sensitivity**.

$$R = \frac{TP}{TP + FN}$$

Accuracy

The **accuracy** is the ratio between the true positives and the true negatives, and the true positives, true negatives, false positives, and false negatives. **Accuracy** alone is a particularly bad quantifier of how well a classifier performs when working with data which is class-imbalanced (i.e. there are not equal numbers of items in each class). In our ground truth dataset, the classes are balanced so measuring accuracy will still be informative.

$$A = \frac{TP + TN}{TP + TN + FP + FN}$$

F1 Score

The  $F_1$  **score** is a weighted average of the **precision** and **recall**. Thus, it takes both false positives and false negatives into account.

$$F1 = \frac{2 * (P * R)}{P + R}$$

Null Accuracy

The **null accuracy** is the accuracy which is obtained by always predicting the most frequent class. Because there are two classes and the tweets within the ground truth dataset equally compose both, the **null accuracy** will always be 50%.

### 2.2.6 *K-Folds Cross-Validation*

Instead of using the entire ground truth dataset as training data for the Naive Bayes classifier, we can split it into a **training set** and a **test set**. The **training set** is fed to the classifier, and the **test set** is used to observe statistics about its performance. Figure 2.2.1 depicts cross-validation using  $k$ -folds, in which we make random splits in the dataset  $k$  times to create several **training set** and **test set** sections. In  $k$ -folds, we then take statistical measurements of how well the classifier performs on the **test set** for each fold. If we made one single split into training and testing sections of our ground truth dataset and only used that single **test set** to gather statistics on the classifier's performance, we would not be able to confirm that those statistics were representative of all the data in the entire ground truth dataset. Instead, we perform 10-fold cross validation, choosing 90% of the data to be the **training set**, and 10% to be the **test set** in each fold. Precision,  $F_1$  score, and recall are calculated within each iteration of the 10 folds, thus utilizing 10 different **test sets**. Finally, we use the averages of those statistics across all folds to measure the overall performance of the classifier.

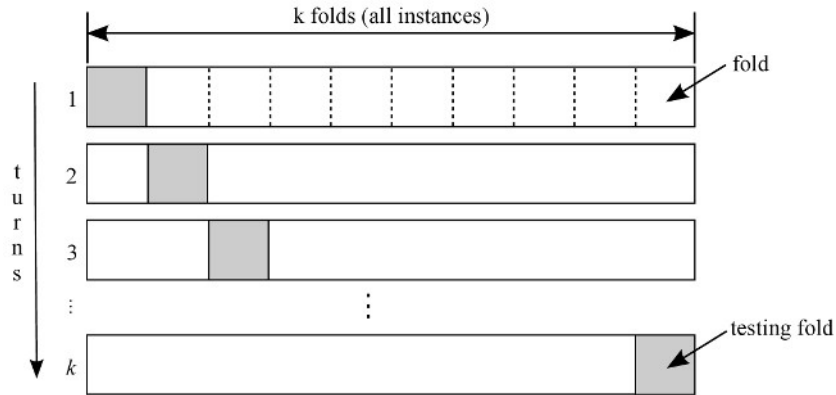


Figure 2.2.1. K-Folds Cross-Validation Example (Borovicka et al., 2012)

### 2.2.7 Training & Testing Naive Bayes

We utilize Scikit Learn’s API for machine learning (Pedregosa et al., 2011) to create a **pipeline**. In Scikit, **pipelines** make managing machine learning algorithms easy by consolidating their parts into one object with configurable attributes.

```

1 sentiment_pipeline = Pipeline([
2     ("vectorizer", TfidfVectorizer(tokenizer=tokenizer.tokenize,
3                                   ngram_range=(1, 3),
4                                   stop_words="english")),
5     ("classifier", MultinomialNB())
6 ])

```

Our **pipeline** contains a vectorizer and a classifier. We change the default arguments for Scikit’s TF-IDF vectorizer to use NLTK’s tweet tokenizer, and specify that we want to calculate our TF-IDF vectors using unigrams, bigrams, and trigrams. Then, we set the vectorizer to use Scikit’s default English language stop words. The Multinomial Naive Bayes Classifier we implement using Scikit includes **Laplace smoothing** ( $\alpha = 1$ ) by default.

Scikit also has a convenient **KFold** object which we utilize to perform cross-validation on our classifier. The goal of  $k$ -folds cross-validation is to train the classifier and acquire statistics on its performance while treating  $k$  different parts of the ground truth dataset as **test sets**. Within a single iteration for 10 iterations, the **test set** will always be 10% of the entire ground truth dataset, however the next iteration will use a different section. We ultimately average those statistics to understand the overall performance of the classifier.

```

1 def confusion_matrices(training_data, num_folds=10):
2     text_training_data = np.array([row[0] for row in training_data])
3     class_training_data = np.array([row[1] for row in training_data])
4     kf = KFold(n_splits=num_folds, random_state=42, shuffle=True)
5     cnf_matrix_test = np.zeros((2, 2), dtype=int)
6     for train_index, test_index in kf.split(text_training_data):
7         text_train, text_test = (text_training_data[train_index],
8                                 text_training_data[test_index])
9         class_train, class_test = (class_training_data[train_index],
10                                  class_training_data[test_index])
11
12         sentiment_pipeline.fit(text_train, class_train)
13         predictions_test = sentiment_pipeline.predict(text_test)
14         cnf_matrix_test += confusion_matrix(class_test, predictions_test)

```

In each iteration of the 10 folds, the above program splits apart a training set and a **test set**.

The classifier is trained on the **training set** using `sentiment_pipeline.fit`, and the classifier's predictions for the **test set** in that fold are used to add toward a confusion matrix which will categorically visualize the performance of the classifier. We also calculate the **precision**, **recall**, and  $F_1$  **score** for each fold's individual **test set**.

## 2.3 Creating the Twitter Bot

We use Tweepy to interact with the Twitter API. It provides a convenient object for streaming Twitter data in real time. The `StreamListener` class can track tweets by searching for specific users, locations, and keywords. For our purposes, it has to be extended to track subtweets.

```

1 class StreamListener(tweepy.StreamListener):
2     def on_status(self, status):
3         id_str = status.id_str
4         screen_name = status.user.screen_name
5         created_at = status.created_at
6         retweeted = status.retweeted
7         in_reply_to = status.in_reply_to_status_id_str
8         text = status.full_text
9
10        # Genericize extra features and clean up the text
11        text = (urls_pattern.sub("GENERIC_URL",
12                                at_mentions_pattern.sub("GENERIC_MENTION",
13                                                        names_pattern.sub("GENERIC_NAME",
14                                                                        text)))
15                .replace("\u2018", "'")
16                .replace("\u2019", "'")
17                .replace("\u201c", "\"")
18                .replace("\u201d", "\"")
19                .replace("&quot;", "\"")
20                .replace("&amp;", "&")
21                .replace("&gt;", ">")
22                .replace("&lt;", "<"))
23
24        tokens = tokenizer.tokenize(text)
25
26        english_tokens = [english_dict.check(token) for token in tokens]

```

```

27     percent_english_words = sum(english_tokens)/float(len(english_tokens))
28
29     # Make sure the tweet contains some English
30     is_english = False
31     if percent_english_words >= 0.1:
32         is_english = True
33
34     # Calculate the probability using the pipeline
35     subtweet_probability = sentiment_pipeline.predict_proba([text]).tolist()[0][1]

```

This part of the program is for live classification of subtweets and gathers information on tweet status objects as they are encountered. We extract data from the tweet object including the ID of that tweet and the text contained within it. We utilize the same techniques for cleaning and genericizing the tweet which we used in preparing our ground truth data for the classifier. The `pipeline` has a `predict_proba` method which takes as its input a list of strings and outputs an array of probabilities for each class. `subtweet_probability`, then, uses that method to predict the probability that the tweet fits the “subtweet” class according to the Naive Bayes classifier and the vectorizer we used in our `pipeline`. We also check that the potential subtweet meets specific requirements before the Twitter bot will interact with it.

```

1  if all([subtweet_probability >= THRESHOLD,
2         "RT" != text[:2], is_english,
3         not retweeted, not in_reply_to]):

```

Included in this conditional statement is a comparison to determine if the probability that the tweet is a subtweet meets a specific threshold. We do not want to call out subtweets unless the probability is high enough. Following this check, we can interact with the tweet in several ways.

```

1  # Quote the tweet
2  api.update_status(("Is this a subtweet? {:.2%} \n" +
3                   "https://twitter.com/{}/status/{}".format(subtweet_probability,
4                                                             screen_name,
5                                                             id_str))
6  # Like the tweet
7  api.create_favorite(id_str)
8
9  # Reply to the tweet
10 api.update_status("@{} Is this a subtweet? {:.2%}".format(screen_name,
11                                                            subtweet_probability),
12                   id_str)

```

To quote the potential subtweet means that the tweet being referenced is embedded within our own tweet with a caption above it. Finally, we instantiate our custom `StreamListener` class and use it to filter through tweets in real time.

```
1 stream_listener = StreamListener()
2 stream = tweepy.Stream(auth=api.auth, listener=stream_listener, tweet_mode="extended")
3
4 stream.filter(follow=user_ids, stall_warnings=True, languages=["en"], async=True)
5 print("Streaming has started.")
6 sleep(DURATION)
7 stream.disconnect()
```

The `user_ids` list contains strings of Twitter user IDs. Every time a user whose ID is in the list sends a tweet, the program will classify that tweet and use the predicted probability that it is a subtweet to call it out on Twitter from the account linked to the application's API credentials. We filter the stream asynchronously in order to use the `sleep` function from the `time` Python library, so we can run the Twitter bot for a limited number of seconds.

# 3

## Results

### 3.1 Ground Truth Dataset

The acquisition of the ground truth dataset relies on assumptions and generalizations of how Twitter users accuse others of subtweeting. For each subtweet and non-subtweet we acquired, we **did not** also acquire every single reply to that tweet, because doing so would violate the API's rate limits (Twitter, 2018). Instead, we maintain that using Twitter's search API, a single reply to a tweet is enough to indicate that it **is** or **is not** a subtweet. By ignoring the associated metadata of tweet objects and their replies, we obtain the unicode text contained within subtweets and non-subtweets for our ground truth dataset.

True Subtweets	True Non-Subtweets
I know who I want to take me home	He’s followed the putrid smell of GENERIC_MENTION which has led him to GENERIC_MENTION’s whereabouts. Some odd neighbor boy was watching him as he approached the house. Liam didn’t hesitate to drink the young one dry. “HONEY, I’M HOME!” He calls, kicking in the door.
One bad chapter doesn’t mean your story is over	YG so low? I guess the only group that sold physicals well is Bigbang GENERIC_URL
What do you offer someone who doesn’t like tea, coffee, or hot chocolate? YOU’RE A GUEST, YOU NEED A HOT DRINK, IT IS THE LAW.	is everything ok?? GENERIC_URL
on the flip side, i agree that “big” accounts have an obligation to do their due diligence to not sic their followers on others	The parade of falsehoods about CIA nominee Gina Haspel GENERIC_URL
im bout to drop like 20lbs and then become a rapper.	Very first day of fortnite got 2 second place games and a third place gg not bad for a trash player :p
It’s hilarious that your only “godly” when it’s convenient for you or when you want to put it on twitter.	Rosie the Corpse Flower bloomed at Tucson Botanical Gardens: See pix allery here GENERIC_URL #Tucson GENERIC_MENTION GENERIC_URL
It really bothers me when people try to take advantage or assume I do shoots for free GENERIC_URL	What is Uncle’s GREATEST prediction?
I just don’t get people lol	Didi Gregious is my new favorite player.
y’all kids today with your “waifu this” and “waifu that,” back in my day our waifus had metal jaws and carried high-frequency blades	GENERIC_MENTION we need to get you drawn as Dekamaster Doggy Krueger
I follow this gorgeous dog account on insta and his owner has such an annoying voice. SHUT UP AND LET ME ENJOY YOUR MAGNIFICENT DOG! I’M NOT HERE FOR YOU!	“Huge Caravan” Of Central Americans Is Headed For The U.S. Border In Hopes Of Asylum GENERIC_URL

When the vectorizer runs over these strings in order to tokenize them and extract features for use in the classifier, GENERIC\_URL, GENERIC\_MENTION, and GENERIC\_NAME stand in place of actual URLs, mentions of user names, and English first names in the original tweets. We use these generic tokens because we want the classifier to probabilistically acknowledge the presence of these features while specifically ignoring the content within them.



### 3.1.1 Characteristics of the Ground Truth Dataset

With tens of thousands of tweets contained within the ground truth dataset, understanding the characteristics of these tweets by example alone is insufficient. To improve our overall understanding of these characteristics, we gathered statistics on the distributions of tweet lengths, punctuating characters, stop words, and unique words in both the **true subtweets** and **true non-subtweets** parts of the dataset.

Twitter serves as a micro-blogging framework which limits the lengths of English text posts to 280 characters. Because Twitter is also a multimedia platform which supports embedding URLs, images, and videos, tweets do not necessarily contain this maximum number of characters. In contrast to the popularity of multimedia tweets, subtweets are characterized as exclusively text-based. Between both **true subtweets** and **true non-subtweets**, Figure 3.1.1 illustrates that the ground truth dataset contains more short tweets ( $1 \leq length \leq 125$ ) than long tweets ( $125 < length \leq 280$ ).




Figure 3.1.1. Histogram of Ground Truth Tweet Lengths  
 $(\mu_s = 103.71, \sigma_s = 71.99), (\mu_n = 105.63, \sigma_n = 74.95)$

Twitter does not enforce any grammar or syntax regulations in tweets, thus punctuating characters are often ignored altogether for use in the stream-of-consciousness style of writing. As shown in Figure 3.1.2, both **true subtweets** and **true non-subtweets** follow this trend in omitting punctuation, on average containing just under 2 punctuating characters (such as quotation marks, periods, commas, and apostrophes).



Figure 3.1.2. Histogram of Ground Truth Tweet Punctuation  
 $(\mu_s = 1.79, \sigma_s = 1.54), (\mu_n = 2.16, \sigma_n = 1.68)$

English stop words are ignored by the vectorizer because they are too common to help the classifier distinguish between the appropriate features for the classes. As shown in Figure 3.1.3, our ground truth dataset contains tweets which are more likely to contain between 5 and 8 unique stop words than they are to contain fewer. In comparison Figure 3.1.4 illustrates the distribution of the number **unique English words** per tweet in the ground truth dataset. As we expect of text-based posts within 280 characters in length, tweets containing fewer unique English words are more common.




Figure 3.1.3. Histogram of Ground Truth Unique English Stop Words  
 $(\mu_s = 7.14, \sigma_s = 1.31), (\mu_n = 7.05, \sigma_n = 1.39)$




Figure 3.1.4. Histogram of Ground Truth English Unique Words  
 $(\mu_s = 16.80, \sigma_s = 10.95), (\mu_n = 16.25, \sigma_n = 11.15)$

## 3.2 Statistics

### 3.2.1 The Confusion Matrix

We used  $k$ -folds cross-validation to test the performance of our Naive Bayes classifier on each **test set** fold of our ground truth dataset. Each individual **test set** was composed of only 10% of the tweets in the entire dataset. By keeping track of the classifier's predictions on this **test set**, we accumulated a **confusion matrix** of **true positives**, **true negatives**, **false positives**, and **false negatives** for all 10 folds. Figure 3.2.1 illustrates these outcomes in terms of raw counts and normalized over the entire ground truth dataset.




Figure 3.2.1. Confusion Matrix on Test Data from Each Fold on the Ground Truth Dataset

We read the confusion matrix by matching rows with columns to see how well the classifier performs on **true subtweets** and **true non-subtweets** in making classifications to produce **predicted subtweets** and **predicted non-subtweets**. These results indicate that the Naive Bayes classifier performs 5% better when classifying true positives than it does when classifying true negatives. In comparison, it is 5% **worse** at classifying false positives than it is at classifying false negatives.

### 3.2.2 Precision, Recall, & $F_1$

The  $k$ -folds cross-validation we utilized to measure the statistical performance of our classifier was averaged across all 10 folds.

	Precision	Recall	$F_1$ Score
Non-Subtweets	0.7357	0.6988	0.7166
Subtweets	0.7132	0.7490	0.7305

The average  $F_1$  scores for non-subtweets and subtweets are respectively 71.66% and 73.05%.

Thus, our classifier performs similarly well on both.

## 3.3 Known Subtweeters

In order to test how the classifier performed on tweets from known subtweeters, we acquired all publicly available tweets from 11 different accounts from which users had previously subtweeted. Figure 3.3.1 shows the distributions of subtweet probabilities produced by the classifier on all these accounts.




Figure 3.3.1. Distribution of Subtweet Probabilities for 11 Known Subtweeters ( $\mu = 0.520, \sigma = 0.127$ )

Of these 11 accounts, the Naive Bayes classifier predicts that out of 26,928 tweets, 15,538 were subtweets (i.e. the predicted subtweet probability was at least 50%). In comparison, if we only accept classifications of at least 75%, then just 924 tweets (3.4%) were predicted to be subtweets. Of these accounts, we selected the users who produced the minimum percentage of subtweets (27.6%), the median percentage of subtweets (58.8%), and the maximum percentage of subtweets (68.4%). Figure 3.3.2 shows the distributions of the subtweet probabilities for all the tweets from these users.




Figure 3.3.2. Distribution of Subtweet Probabilities for 3 Known Subtweeters  
 $(\mu_{min} = 0.430, \sigma_{min} = 0.116), (\mu_{med} = 0.523, \sigma_{med} = 0.128), (\mu_{max} = 0.561, \sigma_{max} = 0.123)$

### 3.4 Most Informative Features

By iterating through the entire vocabulary of features on which the classifier was trained, and sorting those features by the logarithm of the estimated probability that the feature belongs in the subtweet class, we are able to see which features are most informative for classifying tweets as **predicted subtweets**.

Feature	Log Probability
.	-7.655317955479811
,	-8.022304498271094
GENERIC_URL	-8.048671040237405
”	-8.148437784201970
people	-8.427022837390009
?	-8.552068283434405
like	-8.672840985124703
don’t	-8.682709277848670
just	-8.744231174479117
i’m	-8.832921046724476
!	-8.863926249605123
it’s	-9.079518769840222
. GENERIC_URL	-9.100407613290404
:	-9.109587924777987
know	-9.183559311149267
...	-9.204228807382403
you’re	-9.210332638698153
twitter	-9.283961663904783
love	-9.328945234571764
*	-9.404169039956173
friends	-9.418393807734482

The only bigram in the top 20 most informative features is “. GENERIC\_URL,” which can only occur when a URL follows the end of a sentence. In comparison, the rest of these terms are unigrams of which 8 are exclusively punctuating characters (period, comma, quotation, question, exclamation mark, colon, ellipsis, and asterisk). It comes as no surprise that subtweeters talk a lot about Twitter.

### 3.5 The Twitter Bot

After training and testing our classifier, we utilized it in creation of a Twitter bot which interacts with **predicted subtweets** in real time. By limiting the minimum **predicted subtweet** probability threshold to 75%, we prevent the Twitter bot from interacting with users too much. Figures 3.5.1 and 3.5.2 show (censored) examples of these interactions.



Figure 3.5.1. Example of the Twitter Bot Quoting Users' Tweets

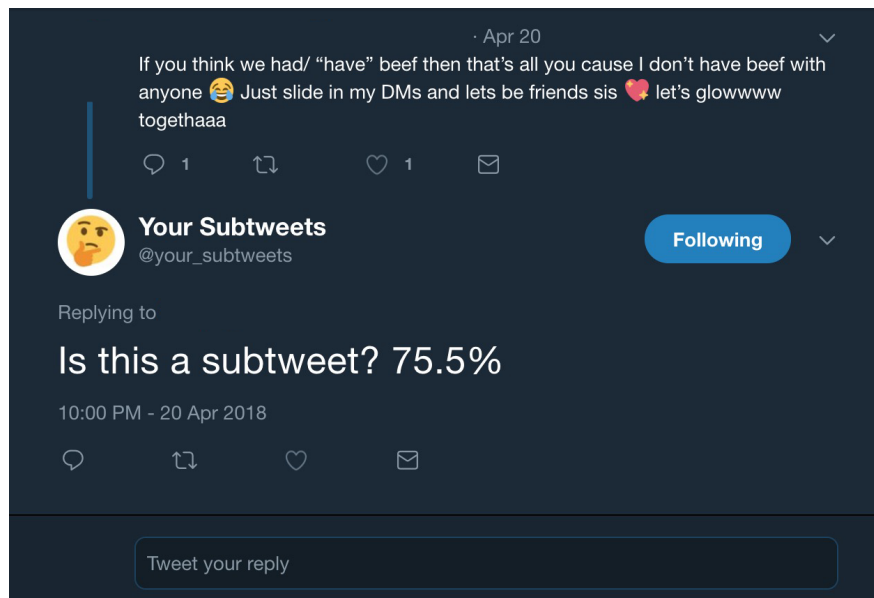


Figure 3.5.2. Example of the Twitter Bot Replying to Users' Tweets

### 3.6 Discussion

We obtained 12,169 **true subtweets** and 21,411 **true non-subtweets** to be used in our ground truth dataset. After confirming they were English and removing duplicates, there were 11,288 **true subtweets** and 19,289 **true non-subtweets** remaining. In order to preserve a balance



between the classes, we limited the size of **true non-subtweets** to be the same as the set of **true subtweets**. Thus, our final ground truth dataset contains 22,576 tweets. Our classifier identifies non-subtweets with an  $F_1$  **score** of 71.66% and subtweets with an  $F_1$  **score** of 73.05%. Subtweets have garnered attention from news organizations (Madrigal, 2014), social scientists (Edwards and Harris, 2016), and governmental officials (Ohlheiser, 2017), but sentiment analysis of subtweets has been entirely unresearched. Utilizing the Twitter API, we acquired data for training and testing a Naive Bayes classifier, and developed a Twitter bot which actively calls-out subtweets in real time.



# 4

## Conclusion

### 4.1 Summary of Project Achievements

We treated a new colloquial form of online harassment—“subtweeting”—as a text classification problem and used the probabilistic Naive Bayes classification algorithm to identify it. We judged the performance of the algorithm and went beyond identification to engage with subtweets, promoting publicity for content which is deliberately written to be unseen by the targeted party. We utilized Twitter’s API to demonstrate a potential use for sentiment analysis on this kind of text.

All data acquired and all programs developed for this project have been made publicly available on GitHub (Noah Segal-Gould, 2018).

### 4.2 Future Work & Considerations

Resources played an important role in this project. With more time, more training data can be acquired. With funds, the Twitter API can be accessed through one of its paid plans to search for tweets from up to 30 days in the past. Because we utilized the free API, the maximum age of an alleged subtweet or non-subtweet when it was acquired could only be one week.

Our implementation of Naive Bayes exclusively classified using features from the unicode text contained within tweets, but other features related to the metadata contained within tweet objects and their replies will probably prove fruitful in producing a better subtweets classifier. We did not test other classification algorithms in favor of pursuing Naive Bayes singularly, but there exists no reason to **not** utilize others. Using topic modeling via methods such as non-negative matrix factorization and latent dirichlet allocation, we can analyze the topics about which most users are tweeting, and potentially group these users into subtweet-topic networked communities.

Finally, we treated this project as a binary classification problem between subtweets and non-subtweets. What would we find if we trained a classifier to distinguish ironic tweets, sarcastic tweets, mocking tweets, and subtweets? What are the features of the figurative language used in each of these examples? These questions motivate further research.

# Bibliography

- Go, Alec, Richa Bhayani, and Lei Huang. 2009. *Twitter Sentiment Classification using Distant Supervision*, CS224N Project Report, Stanford, 12.
- Stone, Biz. 2006. *Introducing the Twitter API*, available at [https://blog.twitter.com/official/en\\_us/a/2006/introducing-the-twitter-api.html](https://blog.twitter.com/official/en_us/a/2006/introducing-the-twitter-api.html).
- Twitter. 2018. *Twitter API Docs*, available at <https://developer.twitter.com/en/docs>.
- . 2016. *Twitter Terms of Service*, available at <https://twitter.com/en/tos>.
- Roesslein, Joshua. 2009. *tweepy Documentation* 5, available at <http://docs.tweepy.org/en/v3.5.0/>.
- Twitter. 2018. *Tweet objects*, available at <https://developer.twitter.com/en/docs/tweets/data-dictionary/overview/tweet-object>.
- Bird, Steven and Edward Loper. 2004. *NLTK: the natural language toolkit*, Proceedings of the ACL 2004 on Interactive poster and demonstration sessions, 31.
- Zhang, Harry. 2004. *The optimality of naive Bayes*, AA 1, no. 2, 3.
- Borovicka, Tomas, Marcel Jirina Jr, Pavel Kordik, and Marcel Jirina. 2012. *Selecting representative data sets*.
- Rae, Chelsea. 2009. *I hate when i see people...*, available at [https://twitter.com/Chelsea\\_x\\_Rae/status/6261479092](https://twitter.com/Chelsea_x_Rae/status/6261479092).
- Urban Dictionary. 2010. *Subtweet*, available at <https://www.urbandictionary.com/define.php?term=subtweet>.
- Edwards, Autumn and Christina J Harris. 2016. *To tweet or subtweet?: Impacts of social networking post directness and valence on interpersonal impressions*, Computers in Human Behavior 63, 304–310.
- Madrigal, Alexis C. 2014. *Behind the machine’s back: How social media users avoid getting turned into big data*, available at <https://goo.gl/h36jxx>.
- Dewey, Caitlin. 2016. *Study confirms what you always knew: People who subtweet are terrible*, available at <https://goo.gl/SeV3mx>.

- Hassler, Chelsea. 2016. *Subtweeting looks terrible on you. (you know who you are.)*, available at <https://goo.gl/NCz27z>.
- Ohlheiser, Abby. 2017. *A running list of all the possible subtweets of President Trump from government Twitter accounts*, available at <https://goo.gl/hFh81R>.
- Twitter. 2012. *Twitter turns six*, available at [https://blog.twitter.com/official/en\\_us/a/2012/twitter-turns-six.html](https://blog.twitter.com/official/en_us/a/2012/twitter-turns-six.html).
- Dorsey, Jack. 2006. *inviting coworkers*, available at <https://twitter.com/jack/status/29>.
- Nassirtoussi, Arman Khadjeh, Saeed Aghabozorgi, Teh Ying Wah, and David Chek Ling Ngo. 2014. *Text mining for market prediction: A systematic review*, Expert Systems with Applications **41**, no. 16, 7653–7670.
- Burnap, Pete, Matthew L Williams, Luke Sloan, Omer Rana, William Housley, Adam Edwards, Vincent Knight, Rob Procter, and Alex Voss. 2014. *Tweeting the terror: modelling the social media reaction to the woolwich terrorist attack*, Social Network Analysis and Mining **4**, no. 1, 206.
- Allport, Gordon W. 1954. *The nature of prejudice*.
- Ghosh, Aniruddha, Guofu Li, Tony Veale, Paolo Rosso, Ekaterina Shutova, John Barnden, and Antonio Reyes. 2015. *Semeval-2015 task 11: Sentiment analysis of figurative language in twitter*, Proceedings of the 9th International Workshop on Semantic Evaluation (SemEval 2015), 470–478.
- Van Hee, Cynthia, Els Lefever, and Vronique Host. 2018. *Semeval-2018 Task 3: Irony detection in English Tweets*, Proceedings of the 12th International Workshop on Semantic Evaluation (SemEval-2018).
- Guilbeault, Douglas and Samuel Woolley. 2016. *How Twitter Bots Are Shaping the Election*, available at <https://goo.gl/ez155d>.
- Bhatnagar, Ranjit. 2018. *Pentametrone*, available at <https://twitter.com/pentametrone>.
- Medhat, Walaa, Ahmed Hassan, and Hoda Korashy. 2014. *Sentiment analysis algorithms and applications: A survey*, Ain Shams Engineering Journal **5**, no. 4, 1093–1113.
- Kelly, Ryan. 2016. *PyEnchant: a spellchecking library for Python*.
- Pedregosa, Fabian, Gaël Varoquaux, Alexandre Gramfort, Vincent Michel, Bertrand Thirion, Olivier Grisel, Mathieu Blondel, Peter Prettenhofer, Ron Weiss, and Vincent Dubourg. 2011. *Scikit-learn: Machine learning in Python*, Journal of machine learning research **12**, no. Oct, 2825–2830.
- Twitter. 2018. *Rate Limiting*, available at <https://developer.twitter.com/en/docs/basics/rate-limiting>.
- Noah Segal-Gould. 2018. *Senior-Project-Subtweets*, available at <https://github.com/segalgouldn/Senior-Project-Subtweets>.

# Appendix A

## Acquiring Ground Truth Data

### A.1 subtweets\_downloader.py

```
1  # coding: utf-8
2
3  ##### Script for downloading a ground truth subtweets dataset
4
5  ##### Import libraries for accessing the API and managing JSON data
6
7  # In[ ]:
8
9
10 import tweepy
11 import json
12
13
14 ##### Load the API credentials
15
16 # In[ ]:
17
18
19 consumer_key, consumer_secret, access_token, access_token_secret = (open("../credentials.txt")
20                                                                     .read().split("\n"))
21
22
23 ##### Authenticate the connection to the API using the credentials
24
25 # In[ ]:
26
27
28 auth = tweepy.OAuthHandler(consumer_key, consumer_secret)
29 auth.set_access_token(access_token, access_token_secret)
30
31
32 ##### Connect to the API
33
34 # In[ ]:
35
36
37 api = tweepy.API(auth, wait_on_rate_limit=True, wait_on_rate_limit_notify=True, compression=True)
38
39
40 ##### Define a function for recursively accessing parent tweets
```

```

41
42 # In[ ]:
43
44
45 def first_tweet(tweet_status_object):
46     try:
47         return first_tweet(api.get_status(tweet_status_object.in_reply_to_status_id_str,
48                                           tweet_mode="extended"))
49     except tweepy.TweepError:
50         return tweet_status_object
51
52
53 #### Define a function for finding tweets with replies that specifically do call them subtweets
54
55 # In[ ]:
56
57
58 def get_subtweets(max_tweets=10000000,
59                  query=("subtweet AND @ since:2018-03-01 exclude:retweets filter:replies")):
60     subtweets_ids_list = []
61     subtweets_list = []
62     i = 0
63     for potential_subtweet_reply in tweepy.Cursor(api.search, lang="en",
64                                                  tweet_mode="extended", q=query).items(max_tweets):
65         i += 1
66         potential_subtweet_original = first_tweet(potential_subtweet_reply)
67         if (not potential_subtweet_original.in_reply_to_status_id_str
68             and potential_subtweet_original.user.lang == "en"):
69             if (potential_subtweet_original.id_str in subtweets_ids_list
70                 or "subtweet" in potential_subtweet_original.full_text
71                 or "Subtweet" in potential_subtweet_original.full_text
72                 or "SUBTWEET" in potential_subtweet_original.full_text):
73                 continue
74             else:
75                 subtweets_ids_list.append(potential_subtweet_original.id_str)
76                 subtweets_list.append({"tweet_data": potential_subtweet_original._json,
77                                       "reply": potential_subtweet_reply._json})
78                 with open("../data/other_data/subtweets.json", "w") as outfile:
79                     json.dump(subtweets_list, outfile, indent=4)
80                 print(("Tweet #{0} was a reply to a subtweet: {1}\n"
81                       .format(i, potential_subtweet_original.full_text.replace("\n", " "))))
82     return subtweets_list
83
84
85 #### Show the results
86
87 # In[ ]:
88
89
90 subtweets_list = get_subtweets()
91 print("Total: {}".format(len(subtweets_list)))

```

## A.2 non\_subtweets\_downloader.py

```

1 # coding: utf-8
2
3 #### Script for downloading a ground truth non-subtweets dataset
4
5 #### Import libraries for accessing the API and managing JSON data
6
7 # In[ ]:
8
9
10 import tweepy
11 import json
12

```



```

13
14 # #### Load the API credentials
15
16 # In[ ]:
17
18
19 consumer_key, consumer_secret, access_token, access_token_secret = (open("../credentials.txt")
20                                                                     .read().split("\n"))
21
22
23 # #### Authenticate the connection to the API using the credentials
24
25 # In[ ]:
26
27
28 auth = tweepy.OAuthHandler(consumer_key, consumer_secret)
29 auth.set_access_token(access_token, access_token_secret)
30
31
32 # #### Connect to the API
33
34 # In[ ]:
35
36
37 api = tweepy.API(auth, wait_on_rate_limit=True, wait_on_rate_limit_notify=True, compression=True)
38
39
40 # #### Define a function for recursively accessing parent tweets
41
42 # In[ ]:
43
44
45 def first_tweet(tweet_status_object):
46     try:
47         return first_tweet(api.get_status(tweet_status_object.in_reply_to_status_id_str,
48                                           tweet_mode="extended"))
49     except tweepy.TweepError:
50         return tweet_status_object
51
52
53 # #### Define a function for finding tweets with replies that specifically do not call them subtweets
54
55 # In[ ]:
56
57
58 def get_non_subtweets(max_tweets=10000000,
59                       query="-subtweet AND @ since:2018-03-01 exclude:retweets filter:replies")):
60     non_subtweets_ids_list = []
61     non_subtweets_list = []
62     i = 0
63     for potential_non_subtweet_reply in tweepy.Cursor(api.search, lang="en",
64                                                       tweet_mode="extended", q=query).items(max_tweets):
65         i += 1
66         potential_non_subtweet_original = first_tweet(potential_non_subtweet_reply)
67         if (not potential_non_subtweet_original.in_reply_to_status_id_str
68             and potential_non_subtweet_original.user.lang == "en"):
69             if (potential_non_subtweet_original.id_str in non_subtweets_ids_list
70                 or "subtweet" in potential_non_subtweet_original.full_text
71                 or "Subtweet" in potential_non_subtweet_original.full_text
72                 or "SUBTWEET" in potential_non_subtweet_original.full_text):
73                 continue
74             else:
75                 non_subtweets_ids_list.append(potential_non_subtweet_original.id_str)
76                 non_subtweets_list.append({"tweet_data": potential_non_subtweet_original._json,
77                                           "reply": potential_non_subtweet_reply._json})
78                 with open("../data/other_data/non_subtweets.json", "w") as outfile:
79                     json.dump(non_subtweets_list, outfile, indent=4)
80                 print(("Tweet #{0} was a reply to a non-subtweet: {1}\n")

```

```
81         .format(i, potential_non_subtweet_original.full_text.replace("\n", " ")))
82     return non_subtweets_list
83
84
85     # #### Show the results
86
87     # In[ ]:
88
89
90     non_subtweets_list = get_non_subtweets()
91     print(len(non_subtweets_list))
```

# Appendix B

## Training & Testing the Classifier

### B.1 classifier\_creator.py

```
1  # coding: utf-8
2
3  # ## Using Scikit-Learn and NLTK to build a Naive Bayes Classifier that identifies subtweets
4
5  # ##### In all tables, assume:
6  # * "GENERIC_URL" represents a single URL
7  # * "GENERIC_MENTION" represents a single mention of a username (e.g. "@noah")
8  # * "GENERIC_NAME" represents a single mention of an English first name
9
10 # ##### Import libraries
11
12 # In[1]:
13
14
15 get_ipython().run_line_magic('matplotlib', 'inline')
16
17
18 # In[2]:
19
20
21 from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
22 from sklearn.feature_extraction.text import TfidfVectorizer
23 from sklearn.feature_extraction import text
24 from sklearn.naive_bayes import MultinomialNB
25 from sklearn.model_selection import KFold
26 from sklearn.pipeline import Pipeline
27 from sklearn.externals import joblib
28 from os.path import basename, splitext
29 from random import choice, sample
30 from nltk.corpus import stopwords
31 from string import punctuation
32 from pprint import pprint
33 from glob import glob
34
35 import matplotlib.pyplot as plt
36 import pandas as pd
37 import numpy as np
38
39 import scipy.stats
40 import itertools
```

```

41 import enchant
42 import nltk
43 import json
44 import re
45
46
47 #### Set up some regex patterns
48
49 In[3]:
50
51
52 urls_pattern = re.compile(r'(?i)\b(?:https?://|www\d{0,3}[.]|'
53     '[a-z0-9.\-]+[.] [a-z]{2,4})/)(?:[^\s('
54     ')<>]|\\((?=[^\s()<>]+|\\((?=[^\s()<>]+\\))'
55     ')*\\))+(?:\\((?=[^\s()<>]+|\\((?=[^\s()<>]'
56     '+\\))*\\)|[^\s!()\\[\\]{};: \\'",.<>?'
57     '\\xab\\xbb\\u201c\\u201d\\u2018\\u2019]))')
58
59
60 In[4]:
61
62
63 at_mentions_pattern = re.compile(r'(?<=^|(?<=[^a-zA-Z0-9-\\.]))@([A-Za-z0-9_]+)')
64
65
66 In[5]:
67
68
69 names = open("../data/other_data/first_names.txt").read().split("\n")
70 names_pattern = re.compile(r'\b(?:{ })\b'.format('|'.join(names)))
71
72
73 #### Prepare English dictionary for language detection
74
75 In[6]:
76
77
78 english_dict = enchant.Dict("en_US")
79
80
81 #### Use NLTK's tokenizer instead of Scikit's
82
83 In[7]:
84
85
86 tokenizer = nltk.casual.TweetTokenizer()
87
88
89 #### Prepare for viewing long text in CSVs and ones with really big and small numbers
90
91 In[8]:
92
93
94 pd.set_option("display.height", 1000)
95 pd.set_option("display.max_rows", 500)
96 pd.set_option("display.max_columns", 500)
97 pd.set_option("display.width", 1000)
98 pd.set_option("max_colwidth", 1000)
99
100
101 In[9]:
102
103
104 pd.options.display.float_format = "{:.4f}".format
105
106
107 #### Load the two data files
108 #### Only use tweets with at least 10% English words

```

```

109 # #### Also, make the mentions of usernames, names, and URLs generic
110
111 # In[10]:
112
113
114 def load_data(filename, threshold=0.1):
115     data = [(urls_pattern.sub("GENERIC_URL",
116         at_mentions_pattern.sub("GENERIC_MENTION",
117         names_pattern.sub("GENERIC_NAME",
118         t["tweet_data"]["full_text"])))
119         .replace("\u2018", "'")
120         .replace("\u2019", "'")
121         .replace("\u201c", "\"")
122         .replace("\u201d", "\"")
123         .replace(""", "\"")
124         .replace("&", "&")
125         .replace(">", ">")
126         .replace("<", "<"))
127     for t in json.load(open(filename)):
128         if t["tweet_data"]["lang"] == "en"
129         and t["reply"]["lang"] == "en"
130         and t["tweet_data"]["user"]["lang"] == "en"
131         and t["reply"]["user"]["lang"] == "en":
132         new_data = []
133         for tweet in data:
134             tokens = tokenizer.tokenize(tweet)
135             english_tokens = [english_dict.check(token) for token in tokens]
136             percent_english_words = sum(english_tokens)/len(english_tokens)
137             if percent_english_words >= threshold:
138                 new_data.append(tweet)
139         return new_data
140
141
142 # In[11]:
143
144
145 subtweets_data = load_data("../data/other_data/subtweets.json")
146
147
148 # In[12]:
149
150
151 non_subtweets_data = load_data("../data/other_data/non_subtweets.json")
152
153
154 # #### Remove tweets which are present in both datasets
155
156 # In[13]:
157
158
159 subtweets_data = [tweet for tweet in subtweets_data
160     if tweet not in non_subtweets_data]
161
162
163 # In[14]:
164
165
166 non_subtweets_data = [tweet for tweet in non_subtweets_data
167     if tweet not in subtweets_data]
168
169
170 # #### Show examples
171
172 # In[15]:
173
174
175 print("Subtweets dataset example:")
176 print(choice(subtweets_data))

```



```

245     ("classifier", MultinomialNB())
246 ]])
247
248
249 # #### K-Folds splits up and separates out 10 training
250 # and test sets from the data, from which the classifier
251 # is trained and the confusion matrix and classification reports are updated
252
253 # In[25]:
254
255
256 def confusion_matrices(training_data, num_folds=10):
257     text_training_data = np.array([row[0] for row in training_data])
258     class_training_data = np.array([row[1] for row in training_data])
259     kf = KFold(n_splits=num_folds, random_state=42, shuffle=True)
260
261     cnf_matrix_test = np.zeros((2, 2), dtype=int)
262     cnf_matrix_train = np.zeros((2, 2), dtype=int)
263
264     test_reports = []
265     train_reports = []
266
267     test accuracies = []
268     train accuracies = []
269     for i, (train_index, test_index) in enumerate(kf.split(text_training_data)):
270
271         text_train, text_test = text_training_data[train_index], text_training_data[test_index]
272         class_train, class_test = class_training_data[train_index], class_training_data[test_index]
273
274         sentiment_pipeline.fit(text_train, class_train)
275
276         predictions_test = sentiment_pipeline.predict(text_test)
277         predictions_train = sentiment_pipeline.predict(text_train)
278
279         cnf_matrix_test += confusion_matrix(class_test, predictions_test)
280         cnf_matrix_train += confusion_matrix(class_train, predictions_train)
281
282         print("Test Data Iteration {}".format(i+1))
283
284         test_report = classification_report(class_test, predictions_test, digits=4)
285         test_reports.append(test_report)
286         print(test_report)
287
288         test_accuracy = accuracy_score(class_test, predictions_test)
289         test accuracies.append(test_accuracy)
290         print("Test Data Accuracy: {:.4f}\n".format(test_accuracy))
291         print("="*53)
292
293         print("Train Data Iteration {}".format(i+1))
294
295         train_report = classification_report(class_train, predictions_train, digits=4)
296         train_reports.append(train_report)
297         print(train_report)
298
299         train_accuracy = accuracy_score(class_train, predictions_train)
300         train accuracies.append(train_accuracy)
301         print("Train Data Accuracy: {:.4f}\n".format(train_accuracy))
302         print("="*53)
303
304     def reports_mean(reports):
305         reports_lists_of_strings = [report.split("\n") for report in reports]
306         reports = [[float(e) for e in report_string[2][16:].split()],
307                    [float(e) for e in report_string[3][16:].split()],
308                    [float(e) for e in report_string[5][16:].split()]]
309         for report_string in reports_lists_of_strings]
310         mean_list = np.mean(np.array(reports), axis=0).tolist()
311         print("          precision    recall  f1-score   support")
312         print()

```

```

313         print("non-subtweet      {0:.4f}    {1:.4f}    {2:.4f}    {3:d}".format(mean_list[0][0],
314                                                         mean_list[0][1],
315                                                         mean_list[0][2],
316                                                         int(mean_list[0][3])))
317     print("      subtweet      {0:.4f}    {1:.4f}    {2:.4f}    {3:d}".format(mean_list[1][0],
318                                                         mean_list[1][1],
319                                                         mean_list[1][2],
320                                                         int(mean_list[1][3])))
321     print()
322     print(" avg / total      {0:.4f}    {1:.4f}    {2:.4f}    {3:d}".format(mean_list[2][0],
323                                                         mean_list[2][1],
324                                                         mean_list[2][2],
325                                                         int(mean_list[2][3])))
326     print()
327     print("="*53)
328
329     print("Test Data Averages Across All Folds:")
330     reports_mean(test_reports)
331
332     print("Train Data Averages Across All Folds:")
333     reports_mean(train_reports)
334
335     return {"Test": cnf_matrix_test, "Train": cnf_matrix_train}
336
337
338 # In[26]:
339
340 cnf_matrices = confusion_matrices(training_data)
341 cnf_matrix_test = cnf_matrices["Test"]
342 cnf_matrix_train = cnf_matrices["Train"]
343
344
345 # #### See the most informative features
346 # [How does "MultinomialNB.coef_" work?](https://stackoverflow.com/a/29915740/6147528)
347
348
349 # In[27]:
350
351
352 def most_informative_features(pipeline, n=15000):
353     vectorizer = pipeline.named_steps["vectorizer"]
354     classifier = pipeline.named_steps["classifier"]
355
356     class_labels = classifier.classes_
357
358     feature_names = vectorizer.get_feature_names()
359
360     top_n_class_1 = sorted(zip(classifier.coef_[0], feature_names))[:n]
361     top_n_class_2 = sorted(zip(classifier.coef_[0], feature_names))[-n:]
362
363     return {class_labels[0]: pd.DataFrame({"Log Probability": [tup[0] for tup in top_n_class_1],
364                                             "Feature": [tup[1] for tup in top_n_class_1]}),
365            class_labels[1]: pd.DataFrame({"Log Probability": [tup[0] for tup in reversed(top_n_class_2)],
366                                             "Feature": [tup[1] for tup in reversed(top_n_class_2)]})}
367
368
369 # In[28]:
370
371
372 most_informative_features_all = most_informative_features(sentiment_pipeline)
373
374
375 # In[29]:
376
377
378 most_informative_features_non_subtweet = most_informative_features_all["non-subtweet"]
379
380

```



```

381 # In[30]:
382
383
384 most_informative_features_subtweet = most_informative_features_all["subtweet"]
385
386
387 # In[31]:
388
389
390 final_features = most_informative_features_non_subtweet.join(most_informative_features_subtweet,
391                                                             lsuffix=" (Non-subtweet)",
392                                                             rsuffix=" (Subtweet)")
393 final_features.to_csv("../data/other_data/most_informative_features.csv")
394 final_features.head(50)
395
396
397 # #### Define function for visualizing confusion matrices
398
399 # In[32]:
400
401
402 def plot_confusion_matrix(cm, classes=["non-subtweet", "subtweet"],
403                           title="Confusion Matrix", cmap=plt.cm.Purples):
404
405     cm_normalized = cm.astype("float") / cm.sum(axis=1)[:, np.newaxis]
406
407     plt.imshow(cm, interpolation="nearest", cmap=cmap)
408     plt.colorbar()
409
410     plt.title(title, size=18)
411
412     tick_marks = np.arange(len(classes))
413     plt.xticks(tick_marks, classes, rotation=45, fontsize=14)
414     plt.yticks(tick_marks, classes, fontsize=14)
415
416     thresh = cm.max() / 2.
417     for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
418         plt.text(j, i, "{} {:.0%}".format(cm[i, j], cm_normalized[i, j]),
419                 horizontalalignment="center", size=16,
420                 color="white" if cm[i, j] > thresh else "black")
421
422     plt.tight_layout()
423
424     plt.ylabel("True label", fontsize=14)
425     plt.xlabel("Predicted Label", fontsize=14)
426
427
428 # #### Show the matrices
429
430 # In[33]:
431
432
433 np.set_printoptions(precision=2)
434
435 plt.figure(figsize=(6, 6))
436 plot_confusion_matrix(cnf_matrix_test, title="Test Data Confusion Matrix")
437
438 plt.figure(figsize=(6, 6))
439 plot_confusion_matrix(cnf_matrix_train, title="Train Data Confusion Matrix")
440
441 plt.show()
442
443
444 # #### Update matplotlib style
445
446 # In[34]:
447
448

```

```

449 plt.style.use("fivethirtyeight")
450
451
452 #### Save the classifier for another time
453
454 # In[35]:
455
456
457 joblib.dump(sentiment_pipeline, "../data/other_data/subtweets_classifier.pkl");
458
459
460 #### Print tests for the classifier
461
462 # In[36]:
463
464
465 def process_tweets_for_testing(filenamees):
466     dataframes = {}
467     for filename in filenamees:
468         username = splitext(basename(filename))[0][:-7]
469         dataframes[username] = {}
470
471         user_df = pd.read_csv(filename).dropna()
472         user_df["Text"] = user_df["Text"].str.replace(urls_pattern, "GENERIC_URL")
473         user_df["Text"] = user_df["Text"].str.replace(at_mentions_pattern, "GENERIC_MENTION")
474         user_df["Text"] = user_df["Text"].str.replace(names_pattern, "GENERIC_NAME")
475         user_df["Text"] = user_df["Text"].str.replace("\u2018", "'")
476         user_df["Text"] = user_df["Text"].str.replace("\u2019", "'")
477         user_df["Text"] = user_df["Text"].str.replace("\u201c", "\"")
478         user_df["Text"] = user_df["Text"].str.replace("\u201d", "\"")
479         user_df["Text"] = user_df["Text"].str.replace(""", "\"")
480         user_df["Text"] = user_df["Text"].str.replace("&", "&")
481         user_df["Text"] = user_df["Text"].str.replace(">", ">")
482         user_df["Text"] = user_df["Text"].str.replace("<", "<")
483
484         predictions = sentiment_pipeline.predict_proba(user_df["Text"])[:, 1].tolist()
485         user_df["SubtweetProbability"] = predictions
486
487         dataframes[username]["all"] = user_df
488
489         scores = user_df[["SubtweetProbability"]].rename(columns={"SubtweetProbability": username})
490
491         dataframes[username]["scores"] = scores
492         dataframes[username]["stats"] = scores.describe()
493
494     return dataframes
495
496
497 #### Load the CSV files
498
499 # In[37]:
500
501
502 filenames = glob("../data/data_for_testing/friends_data/*.csv")
503
504
505 # In[38]:
506
507
508 dataframes = process_tweets_for_testing(filenames)
509
510
511 #### Show a random table
512
513 # In[39]:
514
515
516 chosen_username = choice(list(dataframes.keys()))

```

```

517 dataframes[chosen_username]["all"].sort_values(by="SubtweetProbability", ascending=False).head(5)
518
519
520 # ### Prepare statistics on tweets
521
522 # In[40]:
523
524
525 test_df = pd.concat([df_dict["scores"] for df_dict in dataframes.values()], ignore_index=True)
526
527
528 # In[41]:
529
530
531 test_df_stats = test_df.count()
532
533
534 # ### Total number of tweets per account:
535
536 # In[42]:
537
538
539 test_df_stats
540
541
542 # In[43]:
543
544
545 test_df_over = test_df[test_df >= 0.5]
546
547
548 # In[44]:
549
550
551 test_df_over_stats = test_df_over.count()
552
553
554 # ### Total number of classified subtweets (>= 50%) per account
555
556 # In[45]:
557
558
559 test_df_over_stats
560
561
562 # In[46]:
563
564
565 test_df_combined = pd.concat([test_df_over_stats, test_df_stats], axis=1)
566
567
568 # In[47]:
569
570
571 test_df_combined.columns = ["subtweets", "total_tweets"]
572
573
574 # In[48]:
575
576
577 test_df_combined["percent_subtweets"] = test_df_combined["subtweets"]/test_df_combined["total_tweets"]
578
579
580 # In[49]:
581
582
583 test_df_combined.sort_values(by="percent_subtweets", inplace=True)
584

```

```

585
586 # #### Overall stats
587
588 # In[50]:
589
590
591 test_df_combined
592
593
594 # #### Min sub tweeter
595
596 # In[51]:
597
598
599 min_sub =
    ↪ (test_df_combined.loc[test_df_combined.percent_subtweets==test_df_combined.percent_subtweets.min()]
600     .transpose())
601
602
603 # In[52]:
604
605
606 min_sub
607
608
609 # In[53]:
610
611
612 min_name = min_sub.columns[0]
613
614
615 # #### Median sub tweeter
616
617 # In[54]:
618
619
620 med_sub =
    ↪ (test_df_combined.loc[test_df_combined.percent_subtweets==test_df_combined.percent_subtweets.median()]
621     .transpose())
622
623
624 # In[55]:
625
626
627 med_sub
628
629
630 # In[56]:
631
632
633 med_name = med_sub.columns[0]
634
635
636 # #### Maximum sub tweeter
637
638 # In[57]:
639
640
641 max_sub =
    ↪ (test_df_combined.loc[test_df_combined.percent_subtweets==test_df_combined.percent_subtweets.max()]
642     .transpose())
643
644
645 # In[58]:
646
647
648 max_sub
649

```

```

650
651 # In[59]:
652
653
654 max_name = max_sub.columns[0]
655
656
657 #### Plot a histogram with three random users
658
659 # In[60]:
660
661
662 choices = [dataframes[min_name], dataframes[med_name], dataframes[max_name]]
663 scores = [df_dict["scores"][df_dict["scores"].columns[0]].tolist()
664           for df_dict in choices]
665
666 fig = plt.figure(figsize=(16, 9))
667 ax = fig.add_subplot(111)
668
669 n, bins, patches = ax.hist(scores,
670                             bins="scott",
671                             color=["#FA8072", "#B0A6A4", "#FC1501"],
672                             density=True,
673                             label=["Minimum Subtweeter", "Median Subtweeter", "Maximum Subtweeter"],
674                             alpha=0.75)
675
676 stats = [df_dict["stats"][df_dict["stats"].columns[0]].tolist()
677          for df_dict in choices]
678 for e in stats:
679     print(e[1:3])
680 line_1 = scipy.stats.norm.pdf(bins, stats[0][1], stats[0][2])
681 ax.plot(bins, line_1, "--", color="#FA8072", linewidth=2)
682
683 line_2 = scipy.stats.norm.pdf(bins, stats[1][1], stats[1][2])
684 ax.plot(bins, line_2, "--", color="#B0A6A4", linewidth=2)
685
686 line_3 = scipy.stats.norm.pdf(bins, stats[2][1], stats[2][2])
687 ax.plot(bins, line_3, "--", color="#FC1501", linewidth=2)
688
689 ax.set_xticks([float(x/10) for x in range(11)], minor=False)
690 ax.set_title("Distribution of Subtweet Probabilities In User Accounts", fontsize=18)
691 ax.set_xlabel("Probability That Tweet is a Subtweet", fontsize=18)
692 ax.set_ylabel("Percent of Tweets with That Probability", fontsize=18)
693
694 ax.legend()
695
696 plt.show()
697
698
699 #### Plot a histogram with all of them
700 #### First, get some statistics
701
702 # In[61]:
703
704
705 new_tests_df =
706     ↪ pd.concat([df_dict["scores"].rename(columns={df_dict["scores"].columns[0]: "SubtweetProbability"})
707                for df_dict in dataframes.values()], ignore_index=True)
708
709 new_tests_df_stats = new_tests_df.describe()
710
711 #### Then view them
712
713 # In[62]:
714
715
716 new_tests_df.describe()

```

```

717
718
719 # #### Now plot
720
721 # In[63]:
722
723
724 fig = plt.figure(figsize=(16, 9))
725 ax = fig.add_subplot(111)
726
727 n, bins, patches = ax.hist(new_tests_df["SubtweetProbability"].tolist(),
728                             bins="scott",
729                             color="#6E8B3D",
730                             edgecolor="black",
731                             density=True,
732                             alpha=0.75)
733 print(new_tests_df_stats["SubtweetProbability"][1])
734 print(new_tests_df_stats["SubtweetProbability"][2])
735 line = scipy.stats.norm.pdf(bins, new_tests_df_stats["SubtweetProbability"][1],
736                             new_tests_df_stats["SubtweetProbability"][2])
737
738 ax.plot(bins, line, "--", color="#6E8B3D", linewidth=2)
739
740
741 ax.set_xticks([float(x/10) for x in range(11)], minor=False)
742 ax.set_title("Distribution of Subtweet Probabilities In All User Accounts", fontsize=18)
743 ax.set_xlabel("Probability That Tweet is a Subtweet", fontsize=18)
744 ax.set_ylabel("Percent of Tweets with That Probability", fontsize=18)
745
746 ax.legend()
747
748 plt.show()
749
750
751 # #### Statistics on training data
752
753 # #### Remove generic tokens for these statistics
754
755 # In[64]:
756
757
758 subtweets_data = [tweet[0] for tweet in subtweets_data]
759 non_subtweets_data = [tweet[0] for tweet in non_subtweets_data]
760
761
762 # #### Lengths
763
764 # In[65]:
765
766
767 length_data_subtweets = [len(tweet) for tweet in subtweets_data]
768 length_data_non_subtweets = [len(tweet) for tweet in non_subtweets_data]
769
770
771 # In[66]:
772
773
774 length_data_for_stats_subtweets = pd.DataFrame({"Length": length_data_subtweets,
775                                                  "Tweet": subtweets_data})
776 length_data_for_stats_non_subtweets = pd.DataFrame({"Length": length_data_non_subtweets,
777                                                    "Tweet": non_subtweets_data})
778
779
780 # #### Tweet length statistics
781
782 # In[67]:
783
784

```

```

785 length_data_for_stats_subtweets.describe()
786
787
788 # In[68]:
789
790
791 length_data_for_stats_non_subtweets.describe()
792
793
794 # #### Punctuation
795
796 # In[69]:
797
798
799 punctuation_data_subtweets = [len(set(punctuation).intersection(set(tweet)))
800                               for tweet in subtweets_data]
801 punctuation_data_non_subtweets = [len(set(punctuation).intersection(set(tweet)))
802                                   for tweet in non_subtweets_data]
803
804
805 # In[70]:
806
807
808 punctuation_data_for_stats_subtweets = pd.DataFrame({"Punctuation": punctuation_data_subtweets,
809                                                     "Tweet": subtweets_data})
810 punctuation_data_for_stats_non_subtweets = pd.DataFrame({"Punctuation": punctuation_data_non_subtweets,
811                                                         "Tweet": non_subtweets_data})
812
813
814 # In[71]:
815
816
817 punctuation_data_for_stats_subtweets.describe()
818
819
820 # In[72]:
821
822
823 punctuation_data_for_stats_non_subtweets.describe()
824
825
826 # #### Stop words
827
828 # In[73]:
829
830
831 stop_words_data_subtweets = [len(set(stopwords.words("english")).intersection(set(tweet.lower()))
832                                   for tweet in subtweets_data]
833 stop_words_data_non_subtweets = [len(set(stopwords.words("english")).intersection(set(tweet.lower()))
834                                   for tweet in non_subtweets_data]
835
836
837 # In[74]:
838
839
840 stop_words_data_for_stats_subtweets = pd.DataFrame({"Stop words": stop_words_data_subtweets,
841                                                     "Tweet": subtweets_data})
842 stop_words_data_for_stats_non_subtweets = pd.DataFrame({"Stop words": stop_words_data_non_subtweets,
843                                                         "Tweet": non_subtweets_data})
844
845
846 # #### Tweets stop words statistics
847
848 # In[75]:
849
850
851 stop_words_data_for_stats_subtweets.describe()
852

```

```

853
854 # In[76]:
855
856
857 stop_words_data_for_stats_non_subtweets.describe()
858
859
860 ##### Unique words
861
862 # In[77]:
863
864
865 unique_words_data_subtweets = [sum([english_dict.check(token)
866                                     for token in set(tokenizer.tokenize(tweet))])
867                                for tweet in subtweets_data]
868 unique_words_data_non_subtweets = [sum([english_dict.check(token)
869                                         for token in set(tokenizer.tokenize(tweet))])
870                                    for tweet in non_subtweets_data]
871
872
873 # In[78]:
874
875
876 unique_words_data_for_stats_subtweets = pd.DataFrame({"Unique words": unique_words_data_subtweets,
877                                                         "Tweet": subtweets_data})
878 unique_words_data_for_stats_non_subtweets = pd.DataFrame({"Unique words": unique_words_data_non_subtweets,
879                                                            "Tweet": non_subtweets_data})
880
881
882 ##### Tweets unique words statistics
883
884 # In[79]:
885
886
887 unique_words_data_for_stats_subtweets.describe()
888
889
890 # In[80]:
891
892
893 unique_words_data_for_stats_non_subtweets.describe()
894
895
896 ##### Plot them
897
898 # In[81]:
899
900
901 fig = plt.figure(figsize=(16, 9))
902 ax = fig.add_subplot(111)
903
904 n, bins, patches = ax.hist([length_data_subtweets, length_data_non_subtweets],
905                             bins="sturges",
906                             edgecolor="black",
907                             color=["#B4436C", "#3066BE"],
908                             alpha=0.8)
909 ax.legend(["Subtweets", "Non-Subtweets"], loc="best")
910 ax.set_title("Training Dataset Distribution of Tweet Lengths", fontsize=18)
911 ax.set_xlabel("Tweet Length", fontsize=18);
912 ax.set_ylabel("Number of Tweets with That Length", fontsize=18);
913
914 plt.show()
915
916
917 # In[82]:
918
919
920 fig = plt.figure(figsize=(16, 9))

```



```

921 ax = fig.add_subplot(111)
922
923 n, bins, patches = ax.hist([punctuation_data_subtweets, punctuation_data_non_subtweets],
924                             bins="doane",
925                             edgecolor="black",
926                             color=["#B4436C", "#3066BE"],
927                             alpha=0.8)
928 ax.set_xticks(np.arange(0, 13, step=1))
929 ax.legend(["Subtweets", "Non-Subtweets"], loc="best")
930 ax.set_title("Training Dataset Distribution of Punctuation", fontsize=18)
931 ax.set_xlabel("Punctuating Characters in Tweet", fontsize=18)
932 ax.set_ylabel("Number of Tweets with That Number of Punctuating Characters", fontsize=18)
933
934 plt.show()
935
936
937 # In[83]:
938
939
940 fig = plt.figure(figsize=(16, 9))
941 ax = fig.add_subplot(111)
942
943 n, bins, patches = ax.hist([stop_words_data_subtweets, stop_words_data_non_subtweets],
944                             bins="doane",
945                             edgecolor="black",
946                             color=["#B4436C", "#3066BE"],
947                             alpha=0.8)
948
949 ax.legend(["Subtweets", "Non-Subtweets"], loc="best")
950 ax.set_title("Training Dataset Distribution of Stop Words", fontsize=18)
951 ax.set_xlabel("Stop Words in Tweet", fontsize=18)
952 ax.set_ylabel("Number of Tweets with That Number of Stop Words", fontsize=18)
953
954 plt.show()
955
956
957 # In[84]:
958
959
960 fig = plt.figure(figsize=(16, 9))
961 ax = fig.add_subplot(111)
962
963 n, bins, patches = ax.hist([unique_words_data_subtweets, unique_words_data_non_subtweets],
964                             bins="sturges",
965                             edgecolor="black",
966                             color=["#B4436C", "#3066BE"],
967                             alpha=0.8)
968
969 ax.legend(["Subtweets", "Non-Subtweets"], loc="best")
970 ax.set_title("Training Dataset Distribution of Unique Words", fontsize=18)
971 ax.set_xlabel("Unique English Words in Tweet", fontsize=18)
972 ax.set_ylabel("Number of Tweets with That Number of Unique English Words", fontsize=18)
973
974 plt.show()

```

## B.2 live\_subtweets\_classifier.py

```

1  # coding: utf-8
2
3  ### Script for running a Twitter bot that interacts with subtweets
4
5  #### Import some libraries
6
7  # In[1]:
8
9

```

```

10 from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
11 from sklearn.feature_extraction.text import TfidfVectorizer
12 from sklearn.feature_extraction import text
13 from sklearn.naive_bayes import MultinomialNB
14 from sklearn.pipeline import Pipeline
15 from sklearn.model_selection import KFold
16 from sklearn.externals import joblib
17 from nltk.corpus import stopwords
18 from string import punctuation
19 from pprint import pprint
20 from random import choice
21 from time import sleep
22
23 import pandas as pd
24 import numpy as np
25
26 import itertools
27 import enchant
28 import tweepy
29 import nltk
30 import json
31 import re
32
33
34 #### Prepare the probability threshold for interacting with a potential subtweet and the duration for
35 ↪ which the bot should run
36
37 # In[2]:
38
39 THRESHOLD = 0.75 # 75% positives and higher, only
40 DURATION = 60*15 # 60*60*24*7 # 1 week
41
42
43 #### Set up regular expressions for genericizing extra features
44
45 # In[3]:
46
47
48 urls_pattern = re.compile(r'(?i)\b((?:https?://|www\d{0,3}[.]|[a-z0-9.\-]+|
49     '[.] [a-z]{2,4})/)(?:[^\s()<>]|\\([^\s()<>]+|\\([
50     '[^\s()<>]+\\))*\\))+(?:\\([^\s()<>]+|\\([^\s()<>
51     ']+\\))*\\)|[^\s`!()\\[]{};:\\".,<?\\xab\\xbb\\u201
52     '\\c\\u201d\\u2018\\u2019]))')
53
54
55 # In[4]:
56
57
58 at_mentions_pattern = re.compile(r'(?<=^|(?<=[^a-zA-Z0-9-\\.]))@([A-Za-z0-9_]+)')
59
60
61 # In[5]:
62
63
64 names = open("../data/other_data/first_names.txt").read().split("\n")
65 names_pattern = re.compile(r'\b(?:{})\b'.format('|'.join(names)))
66
67
68 #### Load the classifier pipeline which was previously saved
69
70 # In[6]:
71
72
73 sentiment_pipeline = joblib.load("../data/other_data/subtweets_classifier.pkl")
74
75
76 #### Load the Twitter API credentials

```

```

77
78 # In[7]:
79
80
81 consumer_key, consumer_secret, access_token, access_token_secret = (open("../credentials.txt")
82                                                                     .read()
83                                                                     .split("\n"))
84
85
86 # #### Connect to the API
87
88 # In[8]:
89
90
91 auth = tweepy.OAuthHandler(consumer_key, consumer_secret)
92 auth.set_access_token(access_token, access_token_secret)
93 api = tweepy.API(auth, retry_delay=1, timeout=120, # 2 minutes
94                 compression=True,
95                 wait_on_rate_limit=True, wait_on_rate_limit_notify=True)
96
97
98 # #### Create lists to fill with tweets while the bot is streaming
99
100 # In[9]:
101
102
103 subtweets_live_list = []
104 non_subtweets_live_list = []
105
106
107 # #### Use pyenchant to check if words are English
108
109 # In[10]:
110
111
112 english_dict = enchant.Dict("en_US")
113
114
115 # #### Use NLTK for tokenizing
116
117 # In[11]:
118
119
120 tokenizer = nltk.casual.TweetTokenizer()
121
122
123 # #### Create a custom StreamListener class for use with Tweepy
124
125 # In[12]:
126
127
128 class StreamListener(tweepy.StreamListener):
129     def on_status(self, status):
130         choices = ["retweet", "like", "retweet and like", "reply"]
131
132         id_str = status.id_str
133         screen_name = status.user.screen_name
134         created_at = status.created_at
135         retweeted = status.retweeted
136         in_reply_to = status.in_reply_to_status_id_str
137
138         # The text of the tweet vary based on if it's extended or not
139         if "extended_tweet" in status._json:
140             if "full_text" in status._json["extended_tweet"]:
141                 text = status._json["extended_tweet"]["full_text"]
142             else:
143                 pass # Something else?
144         elif "text" in status._json:

```

```

145     text = status._json["text"]
146
147     # Genericize extra features and clean up the text
148     text = (urls_pattern.sub("GENERIC_URL",
149         at_mentions_pattern.sub("GENERIC_MENTION",
150             names_pattern.sub("GENERIC_NAME",
151                 text)))
152         .replace("\u2018", "'")
153         .replace("\u2019", "'")
154         .replace("\u201c", "\"")
155         .replace("\u201d", "\"")
156         .replace("&quot;", "\"")
157         .replace("&", "&")
158         .replace(">", ">")
159         .replace("<", "<"))
160
161     tokens = tokenizer.tokenize(text)
162
163     english_tokens = [english_dict.check(token) for token in tokens]
164     percent_english_words = sum(english_tokens)/float(len(english_tokens))
165
166     # Make sure the tweet is mostly english
167     is_english = False
168     if percent_english_words >= 0.1:
169         is_english = True
170
171     # Calculate the probability using the pipeline
172     positive_probability = sentiment_pipeline.predict_proba([text]).tolist()[0][1]
173
174     row = {"tweet": text,
175         "screen_name": screen_name,
176         "time": created_at,
177         "subtweet_probability": positive_probability}
178
179     print_list = pd.DataFrame([row]).values.tolist()[0]
180
181     # Only treat it as a subtweet if all conditions are met
182     if all([positive_probability >= THRESHOLD,
183         "RT" != text[:2],
184         is_english, not retweeted, not in_reply_to]):
185         try:
186             decision = choice(choices)
187             if decision == "retweet":
188                 api.update_status(("Is this a subtweet? {:.3%} \n" +
189                     "https://twitter.com/{}/status/{}".format(positive_probability,
190                         screen_name,
191                         id_str))
192                     print("Retweet!")
193
194             elif decision == "like":
195                 api.create_favorite(id_str)
196                 print("Like!")
197
198             elif decision == "retweet and like":
199                 api.update_status(("Is this a subtweet? {:.3%} \n" +
200                     "https://twitter.com/{}/status/{}".format(positive_probability,
201                         screen_name,
202                         id_str))
203                     api.create_favorite(id_str)
204                     print("Retweet and like!")
205
206             elif decision == "reply":
207                 api.update_status("@{} Is this a subtweet? {:.3%}"
208                     .format(screen_name,
209                         positive_probability)),
210                     id_str)
211
212

```

```

213         print("Reply!")
214
215         subtweets_live_list.append(row)
216         subtweets_df = pd.DataFrame(subtweets_live_list).sort_values(by="subtweet_probability",
217                                                                    ascending=False)
218
219         ↪ subtweets_df.to_csv("../data/data_from_testing/live_downloaded_data/subtweets_live_data.csv")
220
221         print(("Subtweet from @{0} (Probability of {1:.3%}):\\n" +
222              "Time: {2}\\n" +
223              "Tweet: {3}\\n" +
224              "Total tweets acquired: {4}\\n").format(print_list[0],
225                                                       print_list[1],
226                                                       print_list[2],
227                                                       print_list[3],
228                                                       (len(subtweets_live_list)
229                                                        + len(non_subtweets_live_list))))
229
230         return row
231     except:
232         print("Unable to interact with tweet!")
233     else:
234         non_subtweets_live_list.append(row)
235         non_subtweets_df = pd.DataFrame(non_subtweets_live_list).sort_values(by="subtweet_probability",
236                                                                              ascending=False)
237
238         ↪ non_subtweets_df.to_csv("../data/data_from_testing/live_downloaded_data/non_subtweets_live_data.csv")
239
240     return row
241
242
243     # #### Create a function for downloading IDs if users I follow who also follow me
244
245     # In[13]:
246
247     def get_mutuals():
248         my_followers = [str(user_id) for ids_list in
249                        tweepy.Cursor(api.followers_ids,
250                                    screen_name="NoahSegalGould").pages()
251                          for user_id in ids_list]
252         my_followeds = [str(user_id) for ids_list in
253                        tweepy.Cursor(api.friends_ids,
254                                    screen_name="NoahSegalGould").pages()
255                          for user_id in ids_list]
256
257         my_mutuals = list(set(my_followers) & set(my_followeds))
258
259         bots = ["890031065057853440", "895685688582180864",
260                "89465860397777152", "970553455709446144",
261                "786489395519983617", "975981192817373184"]
262
263         # Remove known twitter bots
264         my_mutuals = [m for m in my_mutuals if m not in bots]
265
266         with open("../data/other_data/NoahSegalGould_Mutuals_ids.json", "w") as outfile:
267             json.dump(my_mutuals, outfile, sort_keys=True, indent=4)
268
269         return my_mutuals
270
271
272     # #### Create a function for downloading IDs of users
273     # who follow my mutuals who are also followed by my mutuals
274
275     # In[14]:
276
277
278

```

```

279 def get_mutuals_and_mutuals_mutuals_ids(mutuals_threshold=250):
280     my_mutuals = get_mutuals()
281     my_mutuals_mutuals = my_mutuals[:]
282
283     for i, mutual in enumerate(my_mutuals):
284         start_time = time()
285         user = api.get_user(user_id=mutual)
286         name = user.screen_name
287         is_protected = user.protected
288         if not is_protected:
289             mutuals_followers = []
290             followers_cursor = tweepy.Cursor(api.followers_ids, user_id=mutual).items()
291             while True:
292                 try:
293                     mutuals_follower = followers_cursor.next()
294                     mutuals_followers.append(str(mutuals_follower))
295                 except tweepy.TweepError:
296                     sleep(30) # 30 seconds
297                     continue
298                 except StopIteration:
299                     break
300             mutuals_followeds = []
301             followeds_cursor = tweepy.Cursor(api.friends_ids, user_id=mutual).items()
302             while True:
303                 try:
304                     mutuals_followed = followeds_cursor.next()
305                     mutuals_followeds.append(str(mutuals_followed))
306                 except tweepy.TweepError:
307                     sleep(30) # 30 seconds
308                     continue
309                 except StopIteration:
310                     break
311             mutuals_mutuals = list(set(mutuals_followers) & set(mutuals_followeds))
312             print("{} mutuals for mutual {}: {}".format(len(mutuals_mutuals), i+1, name))
313             if len(mutuals_mutuals) <= mutuals_threshold: # Ignore my mutuals if they have a lot of mutuals
314                 my_mutuals_mutuals.extend(mutuals_mutuals)
315             else:
316                 print("\tSkipping: {}".format(name))
317         else:
318             continue
319         end_time = time()
320         with open("../data/other_data/NoahSegalGould_Mutuals_and_Mutuals_Mutuals_ids.json", "w") as
321             ↳ outfile:
322                 json.dump(my_mutuals_mutuals, outfile, sort_keys=True, indent=4)
323                 print("{} seconds for getting the mutuals' IDs of mutual {}: {}".format(
324                     end_time - start_time, i+1, name))
325                 my_mutuals_mutuals = [str(mu) for mu in sorted([int(m) for m in list(set(my_mutuals_mutuals))])]
326                 with open("../data/other_data/NoahSegalGould_Mutuals_and_Mutuals_Mutuals_ids.json", "w") as outfile:
327                     json.dump(my_mutuals_mutuals, outfile, indent=4)
328                 return my_mutuals_mutuals
329
330 # In[15]:
331
332
333 # %%time
334 # my_mutuals_mutuals = get_mutuals_and_mutuals_mutuals_ids()
335
336
337 # #### Load the IDs JSON
338
339 # In[16]:
340
341
342 my_mutuals_mutuals =
343     ↳ json.load(open("../data/other_data/NoahSegalGould_Mutuals_and_Mutuals_Mutuals_ids.json"))
344

```

```
345 # In[17]:
346
347
348 print("Total number of my mutuals and my mutuals' mutuals: {}".format(len(my_mutuals_mutuals)))
349
350
351 # #### Begin streaming
352
353 # In[18]:
354
355
356 stream_listener = StreamListener()
357 stream = tweepy.Stream(auth=api.auth, listener=stream_listener, tweet_mode="extended")
358
359
360 # In[19]:
361
362
363 # %%time
364 # stream.filter(locations=[-73.920176, 42.009637, -73.899739, 42.033421],
365 # stall_warnings=True, languages=["en"], async=True)
366 stream.filter(follow=my_mutuals_mutuals, stall_warnings=True, languages=["en"], async=True)
367 print("Streaming has started.")
368 sleep(DURATION)
369 stream.disconnect()
```